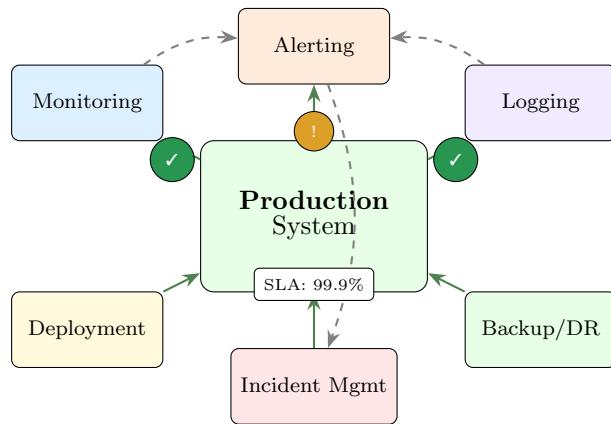


Operational Viewpoint

Architecture Viewpoint Specification

System Operations, Monitoring & Administration



Version: 2.0

Status: Release

Classification: ISO/IEC/IEEE 42010 Compliant

Last Updated: December 12, 2025

Contents

1	Viewpoint Name	3
1.1	Viewpoint Classification	3
1.2	Viewpoint Scope	3
2	Overview	4
2.1	Purpose and Scope	4
2.2	Key Characteristics	4
2.3	Relationship to Other Viewpoints	5
2.4	Operational Architecture Overview	6
3	Concerns	6
3.1	Primary Concerns	6
3.2	Concern-Quality Attribute Mapping	8
4	Anti-Concerns	8
4.1	Out of Scope Topics	8
5	Typical Stakeholders	9
5.1	Primary Stakeholders	10
5.2	Secondary Stakeholders	10
5.3	Stakeholder Concern Matrix	11
6	Model Types	11
6.1	Model Type Catalog	11
6.2	Model Type Relationships	13
7	Model Languages	13
7.1	Observability Diagram Notation	13
7.2	Incident Severity Levels	14
7.3	SLO Specification Format	14
7.4	Runbook Structure	14
7.5	Deployment Pipeline Notation	16
7.6	Tabular Specifications	16
7.6.1	Alert Definition Table	16
7.6.2	Backup Schedule Table	16
8	Viewpoint Metamodels	16
8.1	Core Metamodel	17
8.2	Entity Definitions	17
8.3	Relationship Definitions	23
9	Conforming Notations	23

9.1	ITIL Framework	24
9.2	SRE Practices	24
9.3	Prometheus/Grafana Ecosystem	24
9.4	OpenTelemetry	24
9.5	Operational Tools Comparison	25
10	Model Correspondence Rules	25
10.1	Deployment View Correspondence	25
10.2	Component-and-Connector View Correspondence	26
10.3	Quality Attribute Correspondence	26
11	Operations on Views	26
11.1	Creation Methods	26
11.1.1	View Development Process	27
11.1.2	Common Patterns	28
11.2	Analysis Methods	30
11.2.1	Availability Analysis	30
11.2.2	Error Budget Analysis	30
11.2.3	Incident Metrics Analysis	30
12	Examples	30
12.1	Example 1: Observability Architecture	31
12.2	Example 2: Incident Response Process	31
12.3	Example 3: SLO Dashboard Specification	32
13	Notes	32
13.1	Operational Maturity Model	32
13.2	On-Call Best Practices	33
13.3	Post-Mortem Best Practices	33
13.4	Common Pitfalls	33
14	Sources	33
14.1	Primary References	34
14.2	Supplementary References	34
14.3	Online Resources	34
A	Operational View Checklist	35
B	Glossary	35

1 Viewpoint Name

Viewpoint Identification	
Name:	Operational Viewpoint
Synonyms:	Operations View, Production View, Administration Viewpoint, System Management View, Runtime Operations View, Site Reliability View, Platform Operations View
Identifier:	VP-OPS-001
Version:	2.0

1.1 Viewpoint Classification

The Operational Viewpoint addresses how a system is operated, administered, monitored, and maintained in production environments. While the Views and Beyond approach does not explicitly define an operational viewpoint as a primary style, operational concerns are critical cross-cutting aspects that intersect with deployment, component-and-connector, and quality attribute perspectives. This viewpoint provides essential guidance for running systems reliably in production.

Table 1: Viewpoint Classification Taxonomy

Attribute	Value
Style Family	Cross-Cutting / Quality-Focused
Primary Focus	Production Operations and System Management
Abstraction Level	Operational / Runtime
Temporal Perspective	Runtime and Ongoing Operations
Related Styles	Deployment, C&C, Service-Oriented
IEEE 42010 Category	Operational Viewpoint
Industry Alignment	Site Reliability Engineering (SRE), DevOps, ITIL

1.2 Viewpoint Scope

The Operational Viewpoint encompasses multiple aspects of system operations:

- **Observability:** Monitoring, logging, tracing, and metrics collection to understand system behavior and health in production.
- **Alerting and Incident Management:** Detection of anomalies, alert routing, incident response procedures, and escalation paths.
- **Deployment and Release:** Procedures for deploying, updating, and rolling back system changes safely.

- **Scaling and Capacity:** Mechanisms for handling load variations, capacity planning, and resource management.
- **Backup and Recovery:** Data protection, disaster recovery procedures, and business continuity planning.
- **Maintenance and Administration:** Routine operational tasks, configuration management, and system administration.
- **Service Level Management:** SLAs, SLOs, SLIs, error budgets, and service quality tracking.
- **Security Operations:** Runtime security monitoring, vulnerability management, and incident response.

2 Overview

The Operational Viewpoint provides a comprehensive framework for documenting how a software system is operated, monitored, and maintained throughout its production lifecycle. In modern systems where availability and reliability are critical, this viewpoint is essential for ensuring systems meet their service level commitments.

2.1 Purpose and Scope

The primary purpose of this viewpoint is to establish a clear understanding of operational requirements, procedures, and mechanisms that enable reliable system operation. This includes defining what needs to be monitored, how problems are detected and resolved, how changes are deployed safely, and how the system recovers from failures.

Viewpoint Definition

The Operational Viewpoint defines the mechanisms, procedures, and organizational structures required to operate a system in production. It encompasses observability infrastructure, incident management processes, deployment procedures, capacity management, disaster recovery capabilities, and service level definitions. This viewpoint bridges development activities with production operations.

2.2 Key Characteristics

The Operational Viewpoint exhibits several distinctive characteristics:

Production Focus: This viewpoint is concerned exclusively with the system in production environments, addressing the realities of operating software at scale with real users and data.

Process and Procedure Emphasis: Beyond technical mechanisms, this viewpoint documents the human processes, runbooks, and procedures essential for effective operations.

Continuous Improvement: Operations is iterative; this viewpoint supports ongoing refinement based on incidents, metrics, and changing requirements.

Cross-Functional Integration: Effective operations require coordination between development, operations, security, and business teams. This viewpoint facilitates that integration.

Measurability: Operational concerns are quantified through metrics, SLOs, and SLAs, enabling objective assessment of operational effectiveness.

2.3 Relationship to Other Viewpoints

The Operational Viewpoint connects to other architectural viewpoints in critical ways:

Table 2: Relationships to Other Viewpoints

Viewpoint	Relationship
Deployment	Deployment topology defines what is operated. Infrastructure choices affect operational capabilities.
Component-and-Connector	Runtime components are the units of monitoring. Connector protocols affect observability options.
Information/Data	Data stores require backup and recovery. Data flows affect monitoring points.
Development	Module structure affects deployment units. Build outputs become operational artifacts.
Security	Security monitoring integrates with operational monitoring. Incident response covers security incidents.
Quality Attribute	Availability, performance, and reliability requirements drive operational mechanisms.

2.4 Operational Architecture Overview

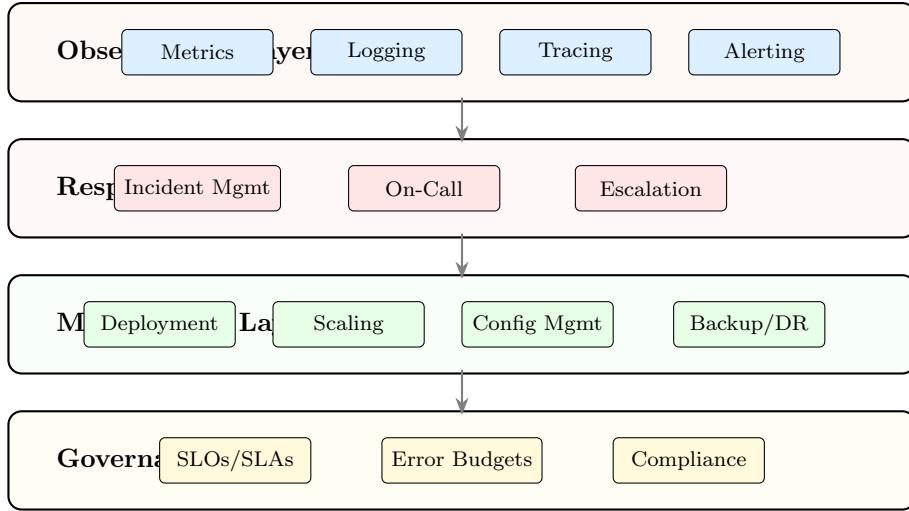


Figure 1: Operational Architecture Layers

3 Concerns

This section enumerates the architectural concerns that the Operational Viewpoint is designed to address. These concerns represent the fundamental questions stakeholders have about system operations.

3.1 Primary Concerns

C1: System Observability

- What metrics indicate system health and performance?
- How are logs collected, aggregated, and searched?
- How are distributed requests traced across services?
- What dashboards provide operational visibility?
- How is observability data retained and accessed?

C2: Alerting and Notification

- What conditions trigger alerts?
- How are alerts routed to appropriate responders?
- What severity levels and escalation paths exist?
- How is alert fatigue minimized?
- What notification channels are used?

C3: Incident Management

- How are incidents detected, classified, and tracked?
- What is the incident response process?

- How are on-call rotations managed?
- What escalation procedures exist?
- How are post-incident reviews conducted?

C4: Deployment and Release

- How are changes deployed to production?
- What deployment strategies are used (blue-green, canary, rolling)?
- How are rollbacks performed?
- What approval and change management processes exist?
- How is deployment frequency and lead time measured?

C5: Scaling and Capacity

- How does the system scale to handle load?
- What auto-scaling policies are configured?
- How is capacity planned and provisioned?
- What are the scaling limits and constraints?
- How is resource utilization optimized?

C6: Backup and Disaster Recovery

- What backup strategies protect data?
- What are the RPO and RTO requirements?
- How is disaster recovery tested?
- What failover mechanisms exist?
- How is business continuity ensured?

C7: Configuration Management

- How is configuration managed across environments?
- How are secrets and credentials handled?
- What configuration drift detection exists?
- How are configuration changes audited?
- What feature flag mechanisms are used?

C8: Service Level Management

- What SLOs define service quality targets?
- How are SLIs measured and reported?
- What SLAs are committed to customers?
- How are error budgets calculated and consumed?
- What happens when SLOs are breached?

C9: Security Operations

- How is security monitored in production?
- What vulnerability management processes exist?

- How are security incidents detected and responded to?
- What compliance monitoring is performed?
- How is access to production controlled?

C10: Maintenance and Administration

- What routine maintenance tasks are required?
- How are maintenance windows scheduled?
- What administrative interfaces exist?
- How is system documentation maintained?
- What automation reduces operational toil?

3.2 Concern-Quality Attribute Mapping

Table 3: Concern to Quality Attribute Mapping

Concern	Availability	Reliability	Performance	Security	Scalability	Maintnc.	Recovertic.	Operability
Observability	●	●	●	○	○	●	○	●
Alerting	●	●	○	○	—	○	○	●
Incident Mgmt	●	●	○	○	—	○	●	●
Deployment	●	○	○	○	○	●	○	●
Scaling	●	○	●	—	●	○	○	○
Backup/DR	●	●	—	○	—	○	●	○
Config Mgmt	○	●	○	●	○	●	○	●
SLO/SLA	●	●	●	○	○	○	○	●
Security Ops	○	○	—	●	—	○	○	○
Maintenance	○	●	○	○	○	●	○	●

● = Primary impact, ○ = Secondary impact, — = Minimal impact

4 Anti-Concerns

Understanding what the Operational Viewpoint is *not* appropriate for helps stakeholders avoid misapplying this viewpoint.

4.1 Out of Scope Topics

AC1: Software Design and Architecture

- Internal component design decisions
- Algorithm implementations
- Code structure and organization
- API design patterns

- Data model design

AC2: Development Process

- Coding standards and practices
- Code review procedures
- Development environment setup
- Unit testing strategies
- Technical debt management

AC3: Business Requirements

- Functional requirement specifications
- User stories and acceptance criteria
- Business process definitions
- Product roadmap planning
- Feature prioritization

AC4: User Experience

- UI/UX design specifications
- User journey mapping
- Accessibility requirements
- Usability testing
- Frontend architecture

AC5: Project Management

- Sprint planning and tracking
- Resource allocation
- Budget management
- Vendor negotiations
- Stakeholder communication plans

Common Misapplications

Avoid using the Operational Viewpoint for:

- Documenting code module structure (use Development Viewpoint)
- Specifying data schemas (use Information Viewpoint)
- Defining infrastructure provisioning (use Deployment Viewpoint)
- Detailing component interactions (use C&C Viewpoint)
- Specifying functional behavior (use Functional Viewpoint)

5 Typical Stakeholders

The Operational Viewpoint serves multiple stakeholder communities with operational responsibilities.

5.1 Primary Stakeholders

Table 4: Primary Stakeholder Analysis

Stakeholder	Role Description	Primary Interests
SRE/Platform Engineers	Ensure system reliability	SLOs, error budgets, automation, incident response, observability
Operations Teams	Manage production systems	Monitoring, alerting, runbooks, maintenance, on-call
DevOps Engineers	Bridge dev and operations	CI/CD, deployment automation, infrastructure as code
Incident Managers	Coordinate incident response	Incident process, escalation, communication, post-mortems
On-Call Engineers	First responders to issues	Alert routing, runbooks, troubleshooting, escalation
Capacity Planners	Ensure adequate resources	Scaling policies, capacity forecasting, cost optimization

5.2 Secondary Stakeholders

Table 5: Secondary Stakeholder Analysis

Stakeholder	Role Description	Primary Interests
Development Teams	Build and maintain software	Deployment process, observability, debugging in production
Security Teams	Protect system assets	Security monitoring, incident response, compliance
Product Managers	Define product direction	SLA commitments, feature flags, release coordination
Customer Support	Assist customers	Incident status, service health, known issues
Executive Leadership	Business accountability	SLA compliance, incident trends, operational costs
Compliance Officers	Ensure regulatory compliance	Audit trails, control effectiveness, compliance reporting

5.3 Stakeholder Concern Matrix

Table 6: Stakeholder-Concern Responsibility Matrix

	<i>Observe</i>	<i>Alerting</i>	<i>Incident</i>	<i>Deploy</i>	<i>Scaling</i>	<i>Backup</i>	<i>Config</i>	<i>SLO</i>	<i>SecOps</i>	<i>Maint.</i>
SRE	R	R	A	A	R	A	A	R	C	A
Ops Team	A	A	R	C	A	R	R	C	C	R
DevOps	C	C	C	R	C	C	R	I	I	C
Incident Mgr	I	C	R	I	I	I	I	A	C	I
On-Call	C	R	R	I	I	C	C	I	C	C
Security	C	C	C	C	I	C	C	I	R	I

R = Responsible, A = Accountable, C = Consulted, I = Informed

6 Model Types

The Operational Viewpoint employs several complementary model types to capture different aspects of system operations.

6.1 Model Type Catalog

MT1: Observability Architecture Diagram

- *Purpose:* Show monitoring, logging, and tracing infrastructure
- *Primary Elements:* Collectors, aggregators, storage, visualization
- *Key Relationships:* Collects-from, forwards-to, stores-in, displays
- *Typical Notation:* Infrastructure diagrams, data flow diagrams

MT2: Alert Configuration Model

- *Purpose:* Document alerting rules, thresholds, and routing
- *Primary Elements:* Alerts, conditions, thresholds, recipients, channels
- *Key Relationships:* Triggers, routes-to, escalates-to
- *Typical Notation:* Alert matrices, routing tables

MT3: Incident Response Process

- *Purpose:* Define incident lifecycle and response procedures
- *Primary Elements:* States, roles, actions, escalation paths
- *Key Relationships:* Transitions-to, escalates-to, notifies
- *Typical Notation:* Process diagrams, state machines, RACI matrices

MT4: Deployment Pipeline Diagram

- *Purpose:* Show deployment stages and promotion paths
- *Primary Elements:* Stages, gates, environments, artifacts

- *Key Relationships:* Promotes-to, requires-approval, deploys-to
- *Typical Notation:* Pipeline diagrams, stage-gate diagrams

MT5: Scaling Policy Model

- *Purpose:* Document auto-scaling rules and capacity management
- *Primary Elements:* Metrics, thresholds, actions, limits
- *Key Relationships:* Triggers, scales, limits
- *Typical Notation:* Policy tables, threshold diagrams

MT6: Disaster Recovery Plan

- *Purpose:* Define backup, recovery, and continuity procedures
- *Primary Elements:* Components, RPO/RTO, procedures, sites
- *Key Relationships:* Backs-up, replicates-to, fails-over-to
- *Typical Notation:* DR topology diagrams, recovery procedures

MT7: SLO/SLA Specification

- *Purpose:* Define service level objectives and agreements
- *Primary Elements:* SLIs, SLOs, SLAs, error budgets
- *Key Relationships:* Measures, targets, commits-to
- *Typical Notation:* SLO tables, error budget charts

MT8: Runbook/Playbook

- *Purpose:* Document operational procedures and troubleshooting
- *Primary Elements:* Steps, decisions, commands, verification
- *Key Relationships:* Follows, branches-to, executes
- *Typical Notation:* Step-by-step procedures, flowcharts

MT9: On-Call Schedule Model

- *Purpose:* Define on-call rotations and coverage
- *Primary Elements:* Teams, schedules, escalation tiers
- *Key Relationships:* Covers, escalates-to, overrides
- *Typical Notation:* Schedule calendars, rotation tables

6.2 Model Type Relationships

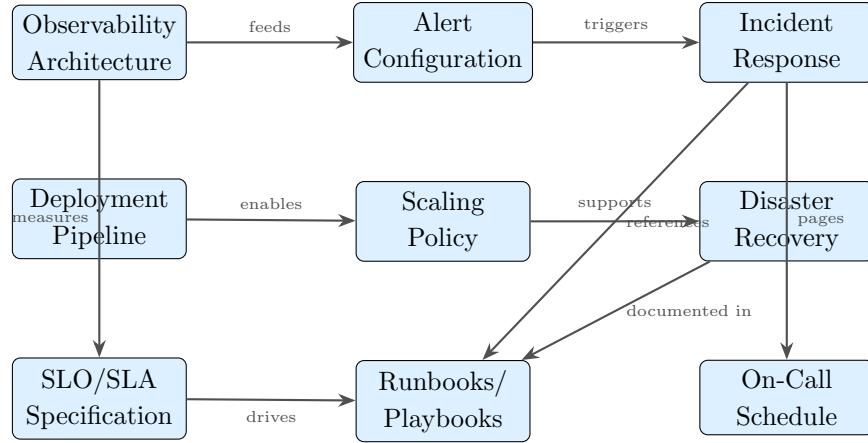


Figure 2: Model Type Dependency Relationships

7 Model Languages

For each model type, specific languages, notations, and techniques are prescribed. This section documents the vocabulary for constructing operational views.

7.1 Observability Diagram Notation

Observability Architecture Notation

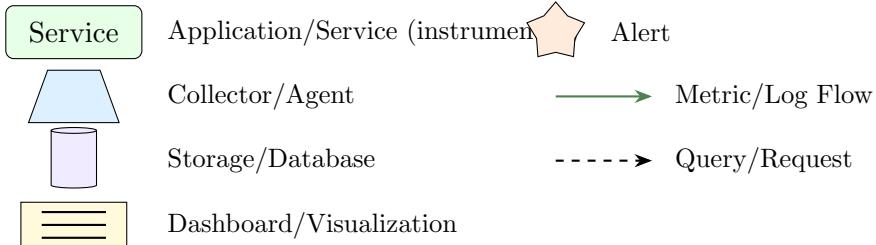


Figure 3: Observability Architecture Notation Legend

7.2 Incident Severity Levels

Table 7: Standard Incident Severity Classification

Level	Name	Criteria	Response
SEV1	Critical	Complete service outage, data loss risk, security breach	Immediate all-hands, exec notification
SEV2	Major	Significant degradation, major feature unavailable	Immediate on-call response, escalation ready
SEV3	Minor	Limited impact, workaround available	Business hours response, next-day acceptable
SEV4	Low	Cosmetic issues, minor inconvenience	Scheduled maintenance window

7.3 SLO Specification Format

SLO Definition Template	
SLO Name:	[Descriptive name for the objective]
Service:	[Service or system this SLO applies to]
SLI (Indicator):	[What is measured - e.g., success rate, latency]
Target:	[Numeric target - e.g., 99.9%, ± 200ms p99]
Window:	[Measurement window - e.g., rolling 30 days]
Error Budget:	[Allowed failures - e.g., 43.2 minutes/month]
Owner:	[Team responsible for this SLO]
Rationale:	[Business justification for target level]

7.4 Runbook Structure

```

1 =====
2 RUNBOOK: [Runbook Title]
3 =====
4
5 METADATA
6 -----
7 Runbook ID:      [RB-XXX]
8 Last Updated:   [Date]
9 Owner:          [Team/Individual]
10 Review Date:    [Next review date]
11 Related Alerts: [Alert names that trigger this runbook]
12
13 OVERVIEW

```

```
14 -----
15 Purpose: [What this runbook addresses]
16 Symptoms: [Observable symptoms that indicate this issue]
17 Impact: [Business/user impact of this issue]
18 Urgency: [SEV level, response time expectations]
19
20 PREREQUISITES
21 -----
22 - Access to [systems/tools required]
23 - Permissions for [actions needed]
24 - Knowledge of [relevant background]
25
26 DIAGNOSTIC STEPS
27 -----
28 1. [First diagnostic action]
29     Expected: [What you should see if this is the issue]
30     Command: [Specific command if applicable]
31
32 2. [Second diagnostic action]
33 ...
34
35 RESOLUTION STEPS
36 -----
37 1. [First resolution action]
38     Verification: [How to verify this step succeeded]
39
40 2. [Second resolution action]
41 ...
42
43 ROLLBACK PROCEDURE
44 -----
45 If resolution fails:
46 1. [Rollback step 1]
47 2. [Rollback step 2]
48
49 ESCALATION
50 -----
51 If unresolved after [time]:
52 - Escalate to: [Team/Individual]
53 - Contact: [Contact information]
54 - Provide: [Information needed for escalation]
55
56 POST-RESOLUTION
57 -----
58 - [ ] Verify service restored
59 - [ ] Update incident ticket
60 - [ ] Notify stakeholders
```

```
61 | - [ ] Schedule post-mortem if needed
62 | =====
```

Listing 1: Runbook Template Structure

7.5 Deployment Pipeline Notation

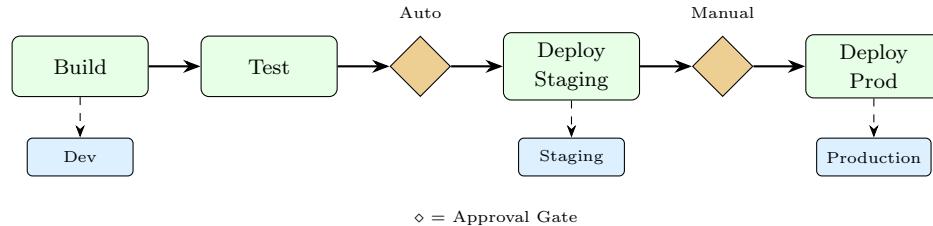


Figure 4: Deployment Pipeline Notation Example

7.6 Tabular Specifications

7.6.1 Alert Definition Table

Table 8: Example Alert Definition Format

Alert	Condition	Threshold	Severity	Team	Runbook
HighErrorRate	error_rate > X	> 1% for 5m	SEV2	Platform	RB-001
HighLatency	p99_latency > X	> 500ms for 10m	SEV3	Platform	RB-002
DiskSpaceLow	disk_used > X	> 85%	SEV3	Infra	RB-003
ServiceDown	health_check	failed 3x	SEV1	On-Call	RB-004

7.6.2 Backup Schedule Table

Table 9: Example Backup Schedule Format

System	Type	Frequency	Retention	RPO	Location
Primary DB	Full	Daily 2AM	30 days	24 hours	S3 us-east-1
Primary DB	Incremental	Hourly	7 days	1 hour	S3 us-east-1
User Files	Full	Weekly Sun	90 days	7 days	S3 us-west-2
Config Store	Full	Daily	14 days	24 hours	S3 us-east-1

8 Viewpoint Metamodels

This section defines the conceptual metamodel underlying the Operational Viewpoint.

8.1 Core Metamodel

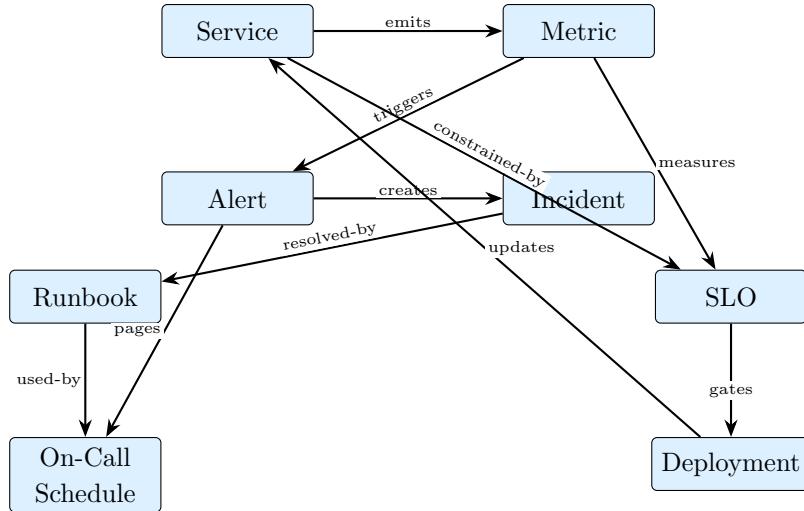


Figure 5: Operational Viewpoint Core Metamodel

8.2 Entity Definitions

Entity: Service

Definition: An operational unit that provides business functionality and is monitored, deployed, and managed as a cohesive entity.

Attributes:

- **serviceId**: Unique identifier
- **name**: Service name
- **description**: Service purpose
- **owner**: Owning team
- **tier**: Criticality tier (Tier 0, 1, 2, 3)
- **dependencies**: Upstream/downstream services
- **healthEndpoint**: Health check URL
- **documentation**: Links to service docs
- **onCallTeam**: Responsible on-call team
- **sloTargets**: Associated SLO references

Constraints:

- Every service must have a defined owner
- Tier 0/1 services must have SLOs defined
- Services must have health check endpoints
- On-call coverage must be assigned

Entity: Metric

Definition: A quantitative measurement collected from a service that indicates its health, performance, or behavior.

Attributes:

- **metricId:** Unique identifier
- **name:** Metric name (following naming conventions)
- **type:** Metric type (counter, gauge, histogram, summary)
- **unit:** Unit of measurement
- **labels:** Dimensional labels for filtering
- **source:** Service or component emitting metric
- **frequency:** Collection frequency
- **retention:** How long metric data is retained
- **aggregations:** Supported aggregation methods

Constraints:

- Metric names must follow naming conventions
- Label cardinality must be bounded
- Critical metrics must have defined retention

Entity: Alert

Definition: A notification triggered when a metric or condition crosses a defined threshold, indicating a potential problem.

Attributes:

- **alertId:** Unique identifier
- **name:** Alert name
- **description:** What the alert indicates
- **condition:** Triggering condition/query
- **threshold:** Threshold value(s)
- **duration:** How long condition must persist
- **severity:** Alert severity level
- **recipients:** Teams/individuals to notify
- **channels:** Notification channels (PagerDuty, Slack, email)
- **runbookLink:** Associated runbook
- **suppressionRules:** When to suppress the alert

Constraints:

- Alerts must have associated runbooks
- SEV1/2 alerts must page on-call
- Alerts should have clear, actionable descriptions
- Suppression rules must be documented

Entity: Incident

Definition: An unplanned event that disrupts or degrades service, requiring response and resolution.

Attributes:

- **incidentId:** Unique identifier
- **title:** Brief description
- **severity:** Incident severity level
- **status:** Current status (open, investigating, mitigating, resolved)
- **affectedServices:** Services impacted
- **startTime:** When incident began
- **detectionTime:** When incident was detected
- **mitigationTime:** When impact was mitigated
- **resolutionTime:** When fully resolved
- **commander:** Incident commander
- **responders:** Response team members
- **timeline:** Sequence of events and actions
- **rootCause:** Identified root cause
- **actionItems:** Follow-up actions

Constraints:

- SEV1/2 incidents require incident commander
- Incidents must have post-mortem within SLA
- Timeline must be maintained during incident
- Action items must be tracked to completion

Entity: SLO (Service Level Objective)

Definition: A target level of service quality measured by a specific indicator, defining acceptable performance boundaries.

Attributes:

- **sloId:** Unique identifier
- **name:** SLO name
- **service:** Associated service
- **sliDefinition:** Service Level Indicator specification
- **target:** Target percentage or value
- **window:** Measurement window (rolling or calendar)
- **errorBudget:** Calculated allowable failures
- **burnRate:** Current consumption rate
- **owner:** Responsible team
- **reviewCadence:** How often SLO is reviewed

Constraints:

- SLO targets must be achievable based on historical data
- SLIs must be measurable and automated
- Error budgets must have defined policies
- SLOs must be reviewed quarterly

Entity: Runbook

Definition: A documented procedure for diagnosing and resolving a specific operational issue or performing a routine task.

Attributes:

- **runbookId:** Unique identifier
- **title:** Descriptive title
- **description:** What the runbook addresses
- **triggeringAlerts:** Alerts that reference this runbook
- **prerequisites:** Required access and knowledge
- **diagnosticSteps:** Investigation procedures
- **resolutionSteps:** Fix procedures
- **rollbackSteps:** How to undo changes
- **escalationPath:** When and how to escalate
- **owner:** Maintaining team
- **lastTested:** When last verified
- **lastUpdated:** Last modification date

Constraints:

- Runbooks must be tested at least quarterly
- Critical runbooks must have multiple reviewers
- Steps must be specific and executable
- Runbooks must be kept current with system changes

Entity: Deployment

Definition: A release of software changes to a production environment, including the process, artifacts, and verification.

Attributes:

- **deploymentId:** Unique identifier
- **service:** Service being deployed
- **version:** Version being deployed
- **previousVersion:** Version being replaced
- **strategy:** Deployment strategy (rolling, blue-green, canary)
- **artifacts:** Deployed artifacts
- **configChanges:** Configuration changes included
- **deployer:** Person/system initiating deployment
- **approvals:** Required approvals obtained
- **startTime:** Deployment start time
- **endTime:** Deployment completion time
- **status:** Current status (pending, in-progress, completed, rolled-back)
- **verificationResults:** Post-deployment checks

Constraints:

- Production deployments require approval
- Deployments must have rollback capability
- Verification checks must pass before completion
- Deployments must be auditable

Entity: On-Call Schedule

Definition: A rotation defining who is responsible for responding to alerts and incidents during specific time periods.

Attributes:

- `scheduleId`: Unique identifier
- `name`: Schedule name
- `team`: Associated team
- `services`: Services covered
- `rotationPeriod`: Length of each rotation (daily, weekly)
- `participants`: Team members in rotation
- `escalationPolicy`: How to escalate if primary unavailable
- `overrides`: Scheduled coverage overrides
- `handoffProcedure`: How shifts are transferred

Constraints:

- 24/7 coverage must be maintained for critical services
- Escalation contacts must always be defined
- On-call load should be distributed fairly
- Handoff procedures must be followed

8.3 Relationship Definitions

Table 10: Metamodel Relationship Definitions

Relationship	Source	Target	Description
emits	Service	Metric	Service produces this metric
triggers	Metric	Alert	Metric condition triggers alert
creates	Alert	Incident	Firing alert creates incident
constrained-by	Service	SLO	Service must meet this SLO
measures	Metric	SLO	Metric is used to calculate SLO
resolved-by	Incident	Runbook	Runbook used to resolve incident
pages	Alert	Schedule	Alert notifies on-call via schedule
updates	Deployment	Service	Deployment changes service
gates	SLO	Deployment	SLO status can block deployment
depends-on	Service	Service	Service requires another service

9 Conforming Notations

Several existing notations and frameworks align with the Operational Viewpoint.

9.1 ITIL Framework

ITIL (Information Technology Infrastructure Library) provides comprehensive service management practices that align with operational concerns.

Conformance Level: Framework alignment for service management processes.

Key Practices: Incident Management, Problem Management, Change Management, Service Level Management.

9.2 SRE Practices

Google's Site Reliability Engineering practices provide detailed guidance for operational excellence.

Conformance Level: Full alignment for reliability engineering.

Key Concepts: SLOs, Error Budgets, Toil Reduction, Blameless Post-mortems.

9.3 Prometheus/Grafana Ecosystem

The Prometheus metrics format and Grafana dashboarding provide de facto standards for observability.

Conformance Level: Tool-specific implementation for monitoring.

Key Formats: PromQL, Grafana JSON dashboard model, Alertmanager configuration.

9.4 OpenTelemetry

OpenTelemetry provides vendor-neutral standards for telemetry data.

Conformance Level: Standard for traces, metrics, and logs.

Key Specifications: OTLP protocol, semantic conventions, SDK specifications.

9.5 Operational Tools Comparison

Table 11: Operational Tool Categories

Category	Tools	Purpose
Metrics	Prometheus, Datadog, New Relic, CloudWatch	Time-series metrics collection and querying
Logging	ELK Stack, Splunk, Loki, CloudWatch Logs	Log aggregation, search, and analysis
Tracing	Jaeger, Zipkin, Datadog APM, X-Ray	Distributed request tracing
Alerting	PagerDuty, OpsGenie, VictorOps	Alert routing and incident management
Dashboarding	Grafana, Datadog, Kibana	Visualization and dashboards
Incident Mgmt	PagerDuty, Jira, ServiceNow	Incident tracking and coordination
On-Call	PagerDuty, OpsGenie, Rootly	On-call scheduling and escalation
Deployment	ArgoCD, Spinnaker, Jenkins, GitHub Actions	CI/CD and deployment automation

10 Model Correspondence Rules

Model correspondence rules define how elements in operational models relate to elements in other architectural views.

10.1 Deployment View Correspondence

Correspondence Rule CR-01: Service to Deployment Mapping

Rule: Every operational service must correspond to deployed artifacts in the deployment view.

Formal Expression:

$$\forall s \in Services_{Ops} : \exists a \in Artifacts_{Deploy} : implements(a, s)$$

Rationale: Ensures operational services have physical deployment presence.

Verification: Deployment manifest review.

Correspondence Rule CR-02: Monitoring Coverage

Rule: Every deployed component must have corresponding observability instrumentation.

Formal Expression:

$$\forall c \in Components_{Deploy} : \exists M \subseteq Metrics : monitors(M, c)$$

Rationale: Ensures complete visibility into deployed systems.

Verification: Metric coverage analysis.

10.2 Component-and-Connector View Correspondence

Correspondence Rule CR-03: Health Check Alignment

Rule: Runtime components must have health endpoints that correspond to operational health checks.

Formal Expression:

$$\forall c \in Components_{C\&C} : \exists h \in HealthChecks : validates(h, c)$$

Rationale: Ensures runtime health can be assessed operationally.

Verification: Health endpoint testing.

10.3 Quality Attribute Correspondence

Correspondence Rule CR-04: SLO to Requirement Mapping

Rule: Every availability/performance quality requirement must have a corresponding SLO.

Formal Expression:

$$\forall qr \in QualityRequirements : type(qr) \in \{availability, performance\} \Rightarrow \exists slo \in SLOs : tracks(slo, qr)$$

Rationale: Ensures quality requirements are operationally measurable.

Verification: Requirements-to-SLO traceability.

11 Operations on Views

This section defines the methods for creating, interpreting, analyzing, and implementing operational views.

11.1 Creation Methods

11.1.1 View Development Process

Step 1: Establish Operational Context

1. Identify services and components requiring operational support
2. Gather availability and reliability requirements
3. Review existing operational processes and tools
4. Identify operational stakeholders and responsibilities
5. Assess current operational maturity level

Step 2: Define Observability Strategy

1. Identify key metrics for each service (RED: Rate, Errors, Duration)
2. Define logging strategy and log levels
3. Plan distributed tracing implementation
4. Select observability tooling stack
5. Define retention policies for telemetry data

Step 3: Establish SLOs

1. Define SLIs based on user-facing experience
2. Set initial SLO targets based on historical performance
3. Calculate error budgets
4. Define error budget policies
5. Establish SLO review cadence

Step 4: Design Alerting Strategy

1. Define alert severity levels and criteria
2. Create alerts based on SLOs (multi-window, multi-burn-rate)
3. Configure alert routing and escalation
4. Implement alert suppression for maintenance
5. Document alert-to-runbook mappings

Step 5: Develop Incident Management Process

1. Define incident severity classification
2. Establish incident response procedures
3. Configure on-call schedules and escalation
4. Create incident communication templates
5. Define post-mortem process

Step 6: Plan Deployment and Recovery

1. Design deployment pipeline with appropriate gates
2. Define rollback procedures
3. Establish backup and recovery procedures
4. Plan disaster recovery and business continuity
5. Document deployment runbooks

Step 7: Create Runbooks

1. Identify scenarios requiring runbooks
2. Document diagnostic and resolution procedures
3. Include escalation paths
4. Test and validate runbooks
5. Establish runbook review cycle

11.1.2 Common Patterns**Pattern: RED Method Metrics**

Context: Need consistent service-level metrics for request-driven services.

Solution: Instrument every service with Rate, Errors, and Duration metrics.

Metrics:

- **Rate:** Requests per second
- **Errors:** Failed requests per second
- **Duration:** Distribution of request latencies

Use When: Monitoring microservices, APIs, web applications.

Pattern: USE Method Metrics

Context: Need consistent resource-level metrics for infrastructure.

Solution: Monitor Utilization, Saturation, and Errors for each resource.

Metrics:

- **Utilization:** Percentage of resource capacity used
- **Saturation:** Degree of queuing/waiting
- **Errors:** Error events for the resource

Use When: Monitoring infrastructure, databases, queues.

Pattern: Multi-Window, Multi-Burn-Rate Alerts

Context: Need alerts that balance responsiveness with noise reduction.

Solution: Use multiple time windows and burn rates to trigger alerts.

Configuration:

- Short window (5m) + high burn rate (14x) = Page immediately
- Medium window (1h) + medium burn rate (6x) = Page
- Long window (6h) + low burn rate (1x) = Ticket

Use When: SLO-based alerting for critical services.

Pattern: Canary Deployment

Context: Need to reduce risk when deploying changes to production.

Solution: Route small percentage of traffic to new version, monitor, then gradually increase.

Process:

1. Deploy new version alongside current version
2. Route 1-5% of traffic to canary
3. Monitor error rates and latency
4. Gradually increase traffic if healthy
5. Auto-rollback if metrics degrade

Use When: Risk-sensitive deployments, services with high traffic.

Table 12: Operational Patterns Summary

Pattern	Description	Use When
RED Metrics	Rate, Errors, Duration for services	Request-driven services
USE Metrics	Utilization, Saturation, Errors	Infrastructure monitoring
Golden Signals	Latency, Traffic, Errors, Saturation	Comprehensive service monitoring
Circuit Breaker	Fail fast when dependency fails	Cascading failure prevention
Bulkhead	Isolate failures to components	Fault isolation
Blue-Green Deploy	Two identical environments	Zero-downtime deployment
Canary Deploy	Gradual traffic shift	Risk-sensitive deployment
Feature Flags	Toggle features without deploy	Progressive rollout
Chaos Engineering	Inject failures deliberately	Resilience validation

11.2 Analysis Methods

11.2.1 Availability Analysis

Availability Calculation

Purpose: Calculate and track service availability.

Formula:

$$\text{Availability} = \frac{\text{TotalTime} - \text{Downtime}}{\text{TotalTime}} \times 100\%$$

Common Targets:

- 99% = 7.3 hours downtime/month
- 99.9% = 43.8 minutes downtime/month
- 99.95% = 21.9 minutes downtime/month
- 99.99% = 4.38 minutes downtime/month

11.2.2 Error Budget Analysis

Error Budget Calculation

Purpose: Track remaining budget for failures/changes.

Formula:

$$\text{ErrorBudget} = (1 - \text{SLOTarget}) \times \text{Window}$$

$$\text{BudgetRemaining} = \text{ErrorBudget} - \text{Consumed}$$

$$\text{BurnRate} = \frac{\text{ErrorRate}}{1 - \text{SLOTarget}}$$

Example: 99.9% SLO over 30 days

- Error Budget = $0.1\% \times 30 \text{ days} = 43.2 \text{ minutes}$
- If 20 minutes consumed, 23.2 minutes remaining
- Burn Rate of 2x means budget exhausted in 15 days

11.2.3 Incident Metrics Analysis

Key Incident Metrics

Purpose: Measure incident management effectiveness.

Key Metrics:

- **MTTD (Mean Time to Detect):** Time from incident start to detection
- **MTTA (Mean Time to Acknowledge):** Time from alert to acknowledgment
- **MTTM (Mean Time to Mitigate):** Time from start to impact mitigation
- **MTTR (Mean Time to Resolve):** Time from start to full resolution
- **Incident Frequency:** Number of incidents per time period
- **Change Failure Rate:** Percentage of deployments causing incidents

12 Examples

12.1 Example 1: Observability Architecture

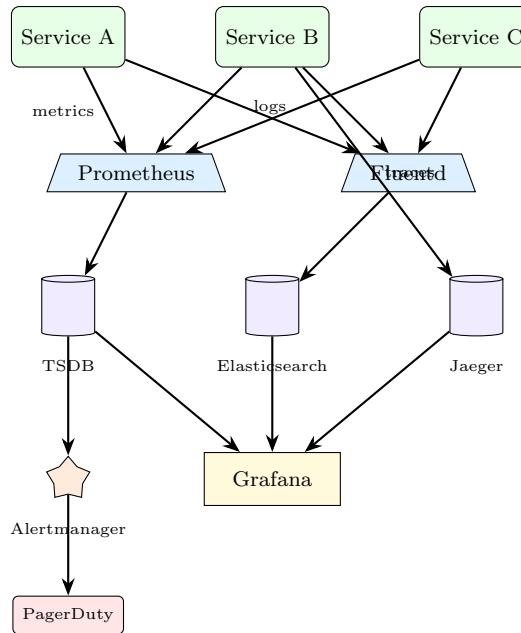


Figure 6: Observability Architecture Example

12.2 Example 2: Incident Response Process

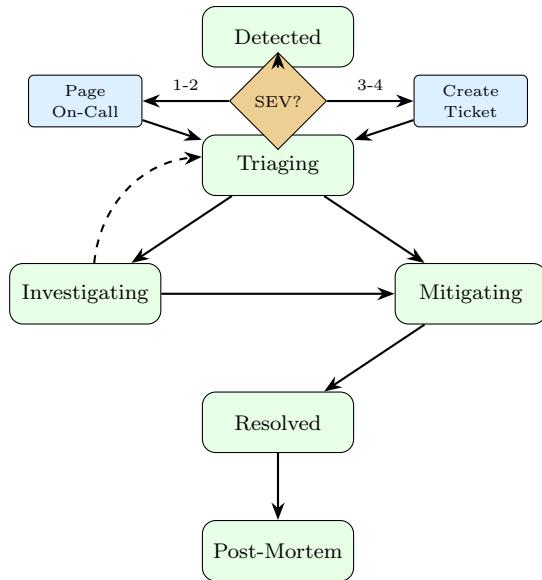


Figure 7: Incident Response State Machine

12.3 Example 3: SLO Dashboard Specification

Table 13: Example SLO Specification

Service: Payment API	
SLO Name	Payment API Availability
SLI	Successful requests / Total requests (excluding 4xx)
Target	99.95%
Window	Rolling 30 days
Error Budget	21.9 minutes/month
Current Status	99.97% (8.7 min consumed, 13.2 min remaining)
Burn Rate	0.4x (healthy)
Owner	Payments Team
SLO Name	Payment API Latency
SLI	Requests completing in 500ms / Total requests
Target	99%
Window	Rolling 30 days
Error Budget	7.3 hours/month
Current Status	99.2% (5.8 hours consumed, 1.5 hours remaining)
Burn Rate	0.8x (caution)
Owner	Payments Team

13 Notes

13.1 Operational Maturity Model

Table 14: Operational Maturity Levels

Level	Name	Characteristics	Capabilities
1	Reactive	Ad-hoc response, manual processes	Basic monitoring, manual deployment
2	Defined	Documented processes, some automation	Defined runbooks, CI/CD basics
3	Measured	SLOs defined, metrics-driven	SLO tracking, automated alerting
4	Managed	Error budgets, proactive management	Predictive scaling, chaos engineering
5	Optimized	Continuous improvement, self-healing	Auto-remediation, ML-based ops

13.2 On-Call Best Practices

On-Call Guidelines

- Limit on-call shifts to reasonable durations (max 1 week)
- Ensure adequate rest between on-call periods
- Provide compensation or time-off for on-call duty
- Maintain clear escalation paths
- Ensure on-call engineers have access to all needed tools
- Conduct regular on-call handoffs
- Track and reduce on-call burden (toil)
- Avoid paging for non-actionable alerts

13.3 Post-Mortem Best Practices

Blameless Post-Mortem Guidelines

- Focus on systems and processes, not individuals
- Assume everyone acted with best intentions and available information
- Document the complete timeline with evidence
- Identify multiple contributing factors
- Generate actionable, assigned, and tracked follow-ups
- Share learnings broadly across the organization
- Review follow-up item completion
- Use post-mortems as learning opportunities

13.4 Common Pitfalls

Common Mistakes to Avoid

1. **Alert Fatigue:** Too many alerts leading to ignored pages
2. **Missing Runbooks:** Alerts without documented response procedures
3. **SLO Without Consequences:** SLOs that don't affect decisions
4. **Untested DR:** Disaster recovery plans never validated
5. **Manual Toil:** Repetitive tasks not automated
6. **Blame Culture:** Post-mortems that assign blame
7. **Incomplete Observability:** Blind spots in monitoring coverage
8. **Overly Complex Deployments:** Deployments without rollback capability

14 Sources

14.1 Primary References

1. Beyer, B., Jones, C., Petoff, J., & Murphy, N. R. (2016). *Site Reliability Engineering: How Google Runs Production Systems*. O'Reilly Media.
2. Beyer, B., Murphy, N. R., Rensin, D. K., Kawahara, K., & Thorne, S. (2018). *The Site Reliability Workbook*. O'Reilly Media.
3. Kim, G., Humble, J., Debois, P., & Willis, J. (2016). *The DevOps Handbook*. IT Revolution Press.
4. Limoncelli, T. A., Chalup, S. R., & Hogan, C. J. (2014). *The Practice of Cloud System Administration*. Addison-Wesley Professional.
5. Majors, C., Fong-Jones, L., & Miranda, G. (2022). *Observability Engineering*. O'Reilly Media.

14.2 Supplementary References

6. Sridharan, C. (2018). *Distributed Systems Observability*. O'Reilly Media.
7. Burns, B. (2018). *Designing Distributed Systems*. O'Reilly Media.
8. Rosenthal, C., & Jones, N. (2020). *Chaos Engineering*. O'Reilly Media.
9. Forsgren, N., Humble, J., & Kim, G. (2018). *Accelerate: The Science of Lean Software and DevOps*. IT Revolution Press.
10. AXELOS. (2019). *ITIL Foundation: ITIL 4 Edition*. TSO.

14.3 Online Resources

- Google SRE Books: <https://sre.google/books/>
- Prometheus Documentation: <https://prometheus.io/docs/>
- OpenTelemetry: <https://opentelemetry.io/>
- PagerDuty Incident Response: <https://response.pagerduty.com/>
- Grafana Documentation: <https://grafana.com/docs/>

A Operational View Checklist

Item	Complete?
Observability	
Metrics collection configured for all services	<input type="checkbox"/>
Logging aggregation in place	<input type="checkbox"/>
Distributed tracing enabled	<input type="checkbox"/>
Dashboards created for key services	<input type="checkbox"/>
Retention policies defined	<input type="checkbox"/>
Alerting	
Alert rules defined based on SLOs	<input type="checkbox"/>
Alert routing configured	<input type="checkbox"/>
Escalation policies in place	<input type="checkbox"/>
All alerts linked to runbooks	<input type="checkbox"/>
Incident Management	
Incident severity levels defined	<input type="checkbox"/>
Incident response process documented	<input type="checkbox"/>
On-call schedules configured	<input type="checkbox"/>
Communication templates ready	<input type="checkbox"/>
Post-mortem process defined	<input type="checkbox"/>
Service Levels	
SLIs defined for critical services	<input type="checkbox"/>
SLO targets established	<input type="checkbox"/>
Error budgets calculated	<input type="checkbox"/>
SLO dashboards created	<input type="checkbox"/>
Deployment & Recovery	
Deployment pipeline documented	<input type="checkbox"/>
Rollback procedures defined	<input type="checkbox"/>
Backup strategy implemented	<input type="checkbox"/>
DR plan documented and tested	<input type="checkbox"/>
Documentation	
Runbooks created for common issues	<input type="checkbox"/>
Service catalog maintained	<input type="checkbox"/>
Architecture diagrams current	<input type="checkbox"/>
On-call documentation complete	<input type="checkbox"/>

B Glossary

Alert	A notification triggered when a monitored condition exceeds a threshold.
Burn Rate	The rate at which error budget is being consumed relative to the SLO window.
Error Budget	The amount of unreliability allowed by an SLO over a given time period.

Incident An unplanned event causing service disruption or degradation.

Mean Time to Detect (MTTD)

Average time from incident start to detection.

Mean Time to Resolve (MTTR)

Average time from incident start to full resolution.

Observability The ability to understand system state from external outputs (metrics, logs, traces).

On-Call A rotation of engineers responsible for responding to production issues.

Post-Mortem A review conducted after an incident to identify causes and improvements.

Recovery Point Objective (RPO)

Maximum acceptable data loss measured in time.

Recovery Time Objective (RTO)

Maximum acceptable time to restore service after failure.

Runbook A documented procedure for handling operational tasks or incidents.

Service Level Agreement (SLA)

A contractual commitment to service quality with consequences.

Service Level Indicator (SLI)

A quantitative measure of service quality.

Service Level Objective (SLO)

A target value for an SLI representing desired service quality.

Toil Repetitive, manual operational work that scales linearly with service growth.