

# Study Plan — OWASP API Security Top 10

## A User Stories Template for Practical Mastery

Compiled for Jordan Suber

## Contents

<b>1 Writing Effective User Stories (Quick Primer)</b>	<b>2</b>
<b>2 Kickoff &amp; Lab Setup</b>	<b>3</b>
<b>3 API1 — Broken Object Level Authorization (BOLA)</b>	<b>4</b>
<b>4 API5 — Broken Function Level Authorization (BFLA)</b>	<b>5</b>
<b>5 API3 — Broken Object Property Level Authorization (BOPLA)</b>	<b>6</b>
<b>6 API2 — Broken Authentication</b>	<b>7</b>
<b>7 API4 — Unrestricted Resource Consumption</b>	<b>8</b>
<b>8 API6 — Unrestricted Access to Sensitive Business Flows</b>	<b>9</b>
<b>9 API7 — Server-Side Request Forgery (SSRF)</b>	<b>10</b>
<b>10 API8 — Security Misconfiguration</b>	<b>11</b>
<b>11 API9 — Improper Inventory Management</b>	<b>12</b>
<b>12 API10 — Unsafe Consumption of APIs</b>	<b>13</b>
<b>13 Cross-Cutting Practices (Do Weekly)</b>	<b>14</b>
<b>14 Capstone Exercise</b>	<b>14</b>

## How to Use This Template

Each card below turns an OWASP API Top 10 (2023) risk into an actionable user story with tasks and verifiable outcomes. Copy any card into your backlog and adjust fields (*Persona*, *Dependencies*, estimates) to your context. Keep stories INVEST-compliant and attach evidence (reports, logs) to each PR/build.

**References.** OWASP API Security Top 10 (2023): API1–API10. See <https://owasp.org/www-project-api-security/> for canonical definitions and guidance.

## 1 Writing Effective User Stories (Quick Primer)

### Required Story Data

- **Epic/Feature** (traceability), **Business Value** (why it matters), **Priority**, **Estimate (SP)**, **Persona**, **Dependencies**, **Assumptions/Risks**.
- **Acceptance Criteria** written in BDD form (Given/When/Then), observable in CI or runtime telemetry.
- **Evidence** links: build artifacts, scanner reports, dashboards, and test logs.

### INVEST Heuristics

- **Independent:** avoids cross-team blocking.
- **Negotiable:** scope can be right-sized.
- **Valuable:** risk reduced, compliance improved, or time saved.
- **Estimable:** bounded by clear AC and DoD.
- **Small:** completes within a sprint (1–5 SP typical).
- **Testable:** pass/fail can be automated.

### Good AC Patterns (Examples)

- **Given** a user with token A and another user's object ID, **When** the GET is attempted, **Then** the API returns 403 and logs an authZ denial with subject/object IDs.
- **Given** login endpoint, **When** 30 requests are sent within 10s from one IP, **Then** rate limiter returns 429 with retry headers and no auth state change.

## **2 Kickoff & Lab Setup**

## W0-Setup — establish a practice repo and CI scaffolding

<b>Epic / Feature</b>	Enablement / Foundations
<b>Business Value</b>	Create a safe space to iterate on risks with automated evidence and reproducibility
<b>Priority / Estimate</b>	Priority: Must SP: 3
<b>Persona</b>	security champion
<b>Dependencies</b>	VCS, CI runner, sample API
<b>Assumptions / Risks</b>	Tooling install time; align languages and package managers across team

**Story** *As a security champion, I want to establish a practice repo and CI scaffolding so that Create a safe space to iterate on risks with automated evidence and reproducibility.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set. • **Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Setup Tasks

- Initialize repo structure: `/openapi`, `/tests`, `/policies`, `/load`, `/docs`.
- Add CI workflow: lint spec (Spectral), run unit tests, publish artifacts and job summary.
- Choose a vulnerable demo API *and* a greenfield API you control.
- Seed `README.md` with architecture diagram, risk register, and Definition of Done (Security).

### Risk-Specific AC for Setup

**Given** CI is configured for the repo

**When** a PR updates tests/spec

**Then** CI emits a summary linking to lint, unit, and coverage reports with pass/fail status

## 3 API1 — Broken Object Level Authorization (BOLA)

## API1-BOLA — enforce object-level authorization checks

<b>Epic / Feature</b>	Access Control
<b>Business Value</b>	Prevent cross-tenant and cross-user data exposure and tampering
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	backend engineer
<b>Dependencies</b>	User/tenant model, auth context in handlers
<b>Assumptions / Risks</b>	Legacy endpoints may bypass middleware; map all object IDs and owners first

**Story** *As a backend engineer, I want to enforce object-level authorization checks so that Prevent cross-tenant and cross-user data exposure and tampering.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set.

**Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed/flagged.

### Tasks

- Inventory resources and ownership (subject→object) per endpoint; record in a matrix.
- Implement server-side ownership checks in controllers/middleware for GET/PUT/PATCH/DELETE.
- Add negative tests that swap IDs in path/body/query to simulate IDOR.
- Log authZ decisions with subject, object, policy, and outcome; redact PII.
- Fail CI if any endpoint lacks an ownership test.

### Risk-Specific AC

**Given** user A and user B exist with distinct object IDs

**When** A requests /objects/{id\_B}

**Then** response is 403 and audit log contains a denied authZ entry with subject/object mapping

## 4 API5 — Broken Function Level Authorization (BFLA)

## API5-BFLA — enforce action/role checks for functions

<b>Epic / Feature</b>	Access Control
<b>Business Value</b>	Stop privilege escalation and unauthorized state changes
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	service owner
<b>Dependencies</b>	RBAC/ABAC policy store, role catalog
<b>Assumptions / Risks</b>	Hidden/legacy admin endpoints; surface and block

**Story** *As a service owner, I want to enforce action/role checks for functions so that Stop privilege escalation and unauthorized state changes.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set.

**Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Tasks

- Build role→action matrix (list/create/update/delete/admin) per endpoint.
- Block undocumented endpoints by default; return 404 or 403 as policy.
- Add tests that attempt admin-only actions with basic user tokens.
- Emit structured authZ logs for each action decision.

### Risk-Specific AC

**Given** a basic role token

**When** calling POST /admin/users

**Then** API returns 403, no side-effects occur, and a denial is recorded

## 5 API3 — Broken Object Property Level Authorization (BO-PLA)

## API3-BOPLA — restrict access to sensitive fields

<b>Epic / Feature</b>	Data Protection
<b>Business Value</b>	Ensure only authorized roles see/modify protected properties
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	API developer
<b>Dependencies</b>	DTO/serializer layer, field visibility rules
<b>Assumptions / Risks</b>	Client filtering is insufficient; enforce on the server

**Story** *As a API developer, I want to restrict access to sensitive fields so that Ensure only authorized roles see/modify protected properties.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set. • **Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Tasks

- Define field visibility per role (e.g., `ssn`, `isAdmin`, secrets).
- Implement server-side projection via DTO/serializers or GraphQL resolvers.
- Add tests asserting sensitive fields never appear for unauthorized roles.
- Create write-guards to block updates to restricted fields.

### Risk-Specific AC

**Given** a basic user token

**When** fetching `/users/me`

**Then** response excludes `ssn` and `isAdmin`; schemas validate absence

## 6 API2 — Broken Authentication

## API2-Auth — harden authentication and sessions

<b>Epic / Feature</b>	Identity & Sessions
<b>Business Value</b>	Reduce account takeover and token abuse by enforcing strong auth flows
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	platform engineer
<b>Dependencies</b>	OAuth2/OIDC or session service, keys/rotation
<b>Assumptions / Risks</b>	Token validation gaps; ensure <code>exp/aud/iss/nbf</code> checks

**Story** *As a platform engineer, I want to harden authentication and sessions so that Reduce account takeover and token abuse by enforcing strong auth flows.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set. • **Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Tasks

- Standardize on OIDC/OAuth2 or signed sessions; rotate signing keys.
- Enforce MFA where appropriate; rate limit login & password reset flows.
- Validate JWT claims and implement secure refresh; prevent replay.
- Add failed-login correlation logs and alerting.

### Risk-Specific AC

**Given** 30 login attempts from one IP within 10s

**When** the attempts are executed

**Then** API returns 429 with retry headers; no session is established

## 7 API4 — Unrestricted Resource Consumption

## API4-URC — apply quotas, limits, and pagination

<b>Epic / Feature</b>	Availability
<b>Business Value</b>	Protect availability and reduce noisy-neighbor effects
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	SRE
<b>Dependencies</b>	Gateway/limiter, pagination patterns
<b>Assumptions / Risks</b>	Large payloads and expensive queries must be controlled

**Story** As a SRE, I want to apply quotas, limits, and pagination so that Protect availability and reduce noisy-neighbor effects.

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set.

**Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed/flagged.

### Tasks

- Define per-endpoint budgets (RPS, burst, concurrency, body size).
- Implement pagination and filters; add 413 on oversize payloads.
- Load-test abusive patterns (k6/Locust); verify graceful 429.
- Dashboards: success/4xx/5xx, rate-limit hits, latency p95/p99.

### Risk-Specific AC

**Given** a client exceeds the configured RPS for /search

**When** requests continue beyond the burst window

**Then** responses are 429 with `Retry-After` and no resource exhaustion occurs

## 8 API6 — Unrestricted Access to Sensitive Business Flows

## API6-Flows — protect high-value flows from automation abuse

<b>Epic / Feature</b>	Abuse Mitigation
<b>Business Value</b>	Prevent fraud/abuse in money/credit/referral flows
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	product security
<b>Dependencies</b>	Risk catalog, friction controls
<b>Assumptions / Risks</b>	Balance friction vs. false positives with telemetry

**Story** *As a product security, I want to protect high-value flows from automation abuse so that Prevent fraud/abuse in money/credit/referral flows.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set. • **Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed/flagged.

### Tasks

- Identify sensitive flows; tag endpoints in OpenAPI with risk annotations.
- Add friction: velocity rules, step-up auth, proof-of-work, or device signals.
- Simulate abuse; measure precision/recall of detections.

### Risk-Specific AC

**Given** 5 coupon redemptions within 60s by same account/IP

**When** the 6th attempt occurs

**Then** flow is blocked or requires step-up auth; event is logged and alerted

## 9 API7 — Server-Side Request Forgery (SSRF)

## API7-SSRF — constrain server egress driven by user input

<b>Epic / Feature</b>	Network Egress
<b>Business Value</b>	Block pivoting to internal services and metadata endpoints
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	platform engineer
<b>Dependencies</b>	Egress proxy, allow-list, DNS re-resolution
<b>Assumptions / Risks</b>	Disable raw URL fetches; sanitize and validate destinations

**Story** *As a platform engineer, I want to constrain server egress driven by user input so that Block pivoting to internal services and metadata endpoints.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set. • **Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed/flagged.

### Tasks

- Route outbound HTTP via proxy; enforce allow-list and DNS re-resolution.
- Block RFC1918/localhost/metadata IP ranges.
- Add tests attempting to reach 169.254.169.254 and internal hosts.
- Log denials with destination details; alert on patterns.

### Risk-Specific AC

**Given** an endpoint accepting a URL

**When** a URL targeting 169.254.169.254 is submitted

**Then** request is denied and logged; no egress connection occurs

## 10 API8 — Security Misconfiguration

## API8-Misconfig — standardize secure defaults and hardening

<b>Epic / Feature</b>	Hardening
<b>Business Value</b>	Reduce attack surface via configs, headers, and container hygiene
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	devops
<b>Dependencies</b>	Baseline headers, pinned images, minimal privileges
<b>Assumptions / Risks</b>	Config drift; codify checks in CI

**Story** *As a devops, I want to standardize secure defaults and hardening so that Reduce attack surface via configs, headers, and container hygiene.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set.

**Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed/flagged.

### Tasks

- Enforce TLS and baseline headers (CORS, HSTS where applicable); minimize error leakage.
- Pin container images; drop root privileges; read-only FS where possible.
- Scan IaC/images (Trivy, tfsec); fail on new high/critical.
- Create a golden service template with secure defaults.

### Risk-Specific AC

**Given** a new service scaffolded from the template

**When** CI runs

**Then** baseline header tests pass and image/IaC scans report no high/critical issues

## 11 API9 — Improper Inventory Management

## API9-Inventory — maintain authoritative API inventory and lifecycle

<b>Epic / Feature</b>	Governance
<b>Business Value</b>	Know every API/version/owner to reduce shadow risk
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	platform owner
<b>Dependencies</b>	API catalog/CMDB, gateway logs
<b>Assumptions / Risks</b>	Rogue endpoints; auto-discovery required

**Story** *As a platform owner, I want to maintain authoritative API inventory and lifecycle so that Know every API/version/owner to reduce shadow risk.*

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set. • **Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Tasks

- Build/extend API catalog with version, owner, data classification, lifecycle.
- Auto-register services from CI on build/release.
- Detect shadow endpoints via gateway logs and repo search.
- Publish deprecation schedules and retirement runbooks.

### Risk-Specific AC

**Given** a new service is released

**When** CI completes

**Then** the service appears in the catalog with owner, version, and classification

## 12 API10 — Unsafe Consumption of APIs

## API10-Consumption — treat upstream APIs as untrusted

<b>Epic / Feature</b>	Third-Party Risk
<b>Business Value</b>	Prevent cascading failures and data misuse from dependencies
<b>Priority / Estimate</b>	Priority: Must SP: 5
<b>Persona</b>	integration engineer
<b>Dependencies</b>	Timeouts/retries/circuit breakers, schema validation
<b>Assumptions / Risks</b>	Provider changes; contract tests mitigate

**Story** As a integration engineer, I want to treat upstream APIs as untrusted so that Prevent cascading failures and data misuse from dependencies.

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set.

**Definition of Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Tasks

- Validate upstream responses against JSON Schema; reject on drift.
- Configure short timeouts, bounded retries, circuit breakers, and fallbacks.
- Add consumer-driven contract tests in CI; pin scopes/permissions.
- Maintain SBOM of integrations and usage terms.

### Risk-Specific AC

**Given** the upstream adds a new required field

**When** contract tests run in CI

**Then** the build fails with a clear schema diff and no production deploy occurs

## 13 Cross-Cutting Practices (Do Weekly)

### Ongoing Tasks

- Update threat models (DFDs, misuse cases) with each change; link to PRs.
- Standardize structured logs: auth decisions, rate-limit hits, SSRF denials.
- Export scanner and test reports (HTML/JSON) as build artifacts.
- Review dashboards (latency p95/p99, % 4xx/5xx) and error budgets.

## 14 Capstone Exercise

## CAP-RedBlue — run a red/blue exercise across API1–API10

<b>Epic / Feature</b>	Readiness
<b>Business Value</b>	Prove detection, response, and hardening in realistic attack flows
<b>Priority / Estimate</b>	Priority: Must SP: 8
<b>Persona</b>	security team
<b>Dependencies</b>	Staging env, attack scripts, runbooks
<b>Assumptions / Risks</b>	Time-box per scenario; record lessons and backlog follow-ups

**Story** As a security team, I want to run a red/blue exercise across API1–API10 so that Prove detection, response, and hardening in realistic attack flows.

**Non-Functional** Security Performance Reliability Privacy Observability

### Acceptance Criteria (Generic, BDD)

**Scenario** Outcome is evidenced in CI

**Given** the target API, spec, and test environment are available

**When** the tasks below are implemented and tests are executed in CI

**Then** the stated outcome is observable in reports, logs, and job summary

**Definition of Ready:** Persona clear; AC drafted; Dependencies known; Estimate set.

**Done:** All ACs pass; Tests green; Security checks; Docs updated; Deployed flagged.

### Tasks

- Simulate BOLA/BOPLA/BFLA/SSRF and business-flow abuse; capture evidence.
- Execute response playbooks (block patterns, rate limits, feature flags).
- Produce a post-mortem with top 5 improvements and owners.

## Checklist Snapshot (Printable)

- API spec linted; security schemes, schemas constrained.
- Ownership/role/field matrices exist with passing tests.
- Rate limits, quotas, and pagination verified under load.
- Egress proxy + allow-list; SSRF denials logged and alerted.
- Golden service template with scans and secure defaults.
- API catalog up to date; deprecation schedules published.
- Third-party contracts tested; timeouts/retries/circuit breakers configured.