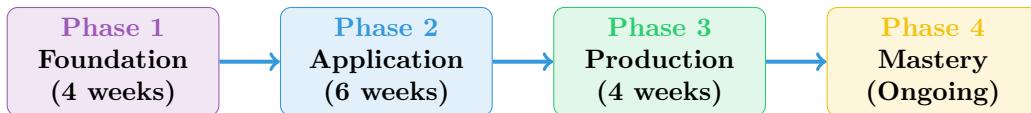


COMPREHENSIVE PROGRAM ROADMAP

Prompt Engineering & LLM Application Development

A Structured Learning Journey from Foundation to Expert-Level Mastery

Technology Agnostic • Provider Independent • Industry Aligned



14–18 Weeks
Core Program

200+ Hours
Learning Time

4 Capstones
Major Projects

50+ Skills
Competencies

Contents

I Program Overview and Navigation	2
1 Executive Roadmap Summary	3
1.1 Program Vision	3
1.2 Program Structure Overview	3
1.3 Learning Outcomes by Phase	3
1.3.1 Phase 1: Foundation (Weeks 1–4)	3
1.3.2 Phase 2: Application Development (Weeks 5–10)	4
1.3.3 Phase 3: Production Engineering (Weeks 11–14)	4
1.3.4 Phase 4: Mastery (Ongoing)	4
1.4 Target Audience Profiles	5
2 Prerequisites and Preparation	6
2.1 Required Knowledge Assessment	6
2.1.1 Technical Skills Matrix	6
2.1.2 Prerequisite Gap Remediation	6
2.2 Environment Setup Checklist	7
2.3 Mindset Preparation	8
2.3.1 Empirical Iteration	8
2.3.2 Provider Agnosticism	8
2.3.3 Continuous Learning	8
2.3.4 Production Orientation	8
II Detailed Learning Path	9
3 Phase 1: Foundation Mastery	10
3.1 Phase 1 Learning Path Visualization	10
3.2 Week 1: Prompt Engineering Fundamentals	10
3.2.1 Learning Objectives	10
3.2.2 Day-by-Day Breakdown	11
3.3 Week 2: Advanced Prompting Patterns	12
3.3.1 Learning Objectives	12
3.3.2 Pattern Deep Dives	13
3.3.3 Week 2 Practical Exercises	14
3.4 Week 3: Systematic Evaluation	15
3.4.1 Learning Objectives	15
3.4.2 Evaluation Framework	15
3.4.3 Week 3 Practical Exercises	16
3.5 Week 4: Specialized Prompting Domains	16

3.5.1	Learning Objectives	16
3.5.2	Code Generation Patterns	17
3.5.3	Structured Output Reliability	17
3.5.4	Multi-Modal Prompting	17
3.6	Phase 1 Capstone: Personal Prompt Cookbook	18
4	Phase 2: Application Development	19
4.1	Phase 2 Learning Path Visualization	19
4.2	Week 5: LLM Application Architecture	19
4.2.1	Learning Objectives	19
4.2.2	Architecture Pattern Catalog	19
4.2.3	Component Design Principles	20
4.3	Week 6: Tool Integration and Function Calling	21
4.3.1	Learning Objectives	21
4.3.2	Tool Design Best Practices	21
4.4	Weeks 7–8: Retrieval-Augmented Generation (RAG)	22
4.4.1	Learning Objectives	22
4.4.2	RAG Pipeline Architecture	23
4.4.3	Chunking Strategies	23
4.4.4	Retrieval Optimization	23
4.5	Week 9: Conversational Agents and Memory	24
4.5.1	Learning Objectives	24
4.5.2	Memory Architecture Comparison	24
4.5.3	Persona Consistency Techniques	25
4.6	Week 10: Integration and Capstone Projects	25
4.6.1	Portfolio Project Requirements	25
5	Phase 3: Production Engineering	27
5.1	Phase 3 Learning Path Visualization	27
5.2	Week 11: System Design for Production	27
5.2.1	Learning Objectives	27
5.2.2	Decision Framework: Prompting vs. Fine-Tuning vs. RAG	27
5.2.3	Cost Modeling Framework	28
5.3	Week 12: Evaluation at Scale	28
5.3.1	Learning Objectives	28
5.3.2	Multi-Level Evaluation Framework	29
5.3.3	LLM-as-Judge Best Practices	29
5.4	Week 13: LLMOps and Production Operations	29
5.4.1	Learning Objectives	29
5.4.2	Observability Stack Components	30
5.4.3	Prompt Version Management	30
5.5	Week 14: Safety, Security, and Ethics	30
5.5.1	Learning Objectives	30
5.5.2	Security Vulnerability Catalog	30
5.5.3	Defense-in-Depth Strategy	31
5.6	Phase 3 Capstone: Production Playbook	32
6	Phase 4: Continuous Mastery	34
6.1	Mastery Development Framework	34
6.1.1	Advanced Technique Areas	34
6.1.2	Specialization Tracks	34
6.1.3	Continuous Learning Practices	35

6.1.4	Community Engagement	35
III	Supporting Materials	36
7	Competency Framework	37
7.1	Skill Progression Model	37
7.2	Detailed Competency Matrix	37
7.2.1	Prompt Engineering Competencies	37
7.2.2	Application Development Competencies	37
7.2.3	Production Engineering Competencies	38
8	Self-Assessment Tools	39
8.1	Phase Readiness Checklists	39
8.1.1	Ready for Phase 2?	39
8.1.2	Ready for Phase 3?	40
8.1.3	Ready for Phase 4 (Mastery)?	41
9	Resources and References	42
9.1	Core Reading List	42
9.1.1	Primary Textbooks	42
9.1.2	Official Documentation (Provider-Agnostic Approach)	42
9.2	Frameworks and Tools	42
9.3	Key Research Papers	43
10	Schedule Options and Pacing Guides	44
10.1	Standard Track (14 Weeks)	44
10.2	Accelerated Track (10 Weeks)	44
10.3	Extended Track (18+ Weeks)	45
10.4	Self-Paced Guidelines	45
A	Quick Reference Cards	46
A.1	Prompting Pattern Quick Reference	46
A.2	Architecture Decision Quick Reference	46
A.3	Security Checklist	47
B	Glossary	48

Part I

Program Overview and Navigation

Chapter 1


Executive Roadmap Summary

1.1 Program Vision

This comprehensive roadmap provides a structured, technology-agnostic pathway for mastering prompt engineering and building production-ready LLM-powered applications. The program is designed around the principle that effective LLM engineering requires a combination of foundational prompt crafting skills, software engineering best practices, and operational excellence.

The roadmap emphasizes practical application over theoretical knowledge, with each phase building upon the previous one through hands-on projects, iterative skill development, and progressive complexity. Learners emerge not just as prompt engineers, but as complete AI application developers capable of designing, building, deploying, and maintaining sophisticated LLM systems.

1.2 Program Structure Overview



1.3 Learning Outcomes by Phase

1.3.1 Phase 1: Foundation (Weeks 1–4)

Upon completing Phase 1, learners will demonstrate mastery of core prompting principles applicable across any LLM platform:

- **Prompt Construction:** Design clear, effective prompts that consistently produce high-quality outputs regardless of the underlying model

- **Pattern Recognition:** Identify and apply 25+ distinct prompting patterns to appropriate use cases
- **Reasoning Enhancement:** Implement chain-of-thought, self-consistency, and other reasoning amplification techniques
- **Quality Evaluation:** Systematically assess and iterate on prompt performance using both qualitative and quantitative methods
- **Multi-Modal Competence:** Extend prompting skills to code generation, structured outputs, and vision-language tasks

1.3.2 Phase 2: Application Development (Weeks 5–10)

Phase 2 transforms prompt engineers into application developers:

- **Architecture Design:** Select and implement appropriate architectural patterns for LLM applications
- **Tool Integration:** Enable LLMs to interact with external systems through function calling and tool use
- **Knowledge Augmentation:** Build retrieval-augmented generation (RAG) systems for knowledge-intensive applications
- **Conversational Systems:** Design stateful agents with memory, context management, and persona consistency
- **Provider Abstraction:** Create applications that work across multiple LLM providers without code changes

1.3.3 Phase 3: Production Engineering (Weeks 11–14)

Phase 3 focuses on operational excellence and production readiness:

- **System Design:** Make informed trade-offs between prompting, fine-tuning, and RAG approaches
- **Quality Assurance:** Implement comprehensive evaluation pipelines with automated and human-in-the-loop assessment
- **Observability:** Deploy monitoring, logging, and alerting systems for LLM applications
- **Cost Optimization:** Manage inference costs through caching, model selection, and prompt optimization
- **Security Hardening:** Protect against prompt injection, data leakage, and other LLM-specific vulnerabilities

1.3.4 Phase 4: Mastery (Ongoing)

Phase 4 represents continuous professional growth:

- **Advanced Techniques:** Master emerging patterns like constitutional AI, recursive self-improvement, and multi-agent orchestration

- **Research Integration:** Stay current with academic developments and integrate new techniques into practice
- **Domain Specialization:** Develop deep expertise in specific verticals (healthcare, finance, legal, etc.)
- **Technical Leadership:** Guide teams and organizations in LLM adoption and best practices

1.4 Target Audience Profiles

This roadmap accommodates multiple learner profiles with tailored pathways:

Profile	Background & Goals	Recommended Track
Software Engineer	Experienced developer seeking to integrate LLM capabilities into existing applications	Standard (14 weeks)
Data Scientist	ML practitioner transitioning to LLM-focused roles; familiar with model evaluation	Accelerated (10 weeks)
Product Manager	Technical PM requiring deep understanding of LLM capabilities and limitations	Foundation+ (8 weeks)
Technical Writer	Content creator using generative AI for documentation and communication	Custom (Phase 1 + select modules)
Career Changer	Professional from another field entering AI/ML; strong analytical skills	Extended (18+ weeks)

Chapter 2

Prerequisites and Preparation

2.1 Required Knowledge Assessment

Before beginning the program, learners should honestly assess their current competencies against the following requirements. This self-assessment determines your recommended starting point and pace.

⚠ Prerequisites

Essential Prerequisites (Must Have):

- Proficiency in at least one programming language (Python strongly recommended)
- Understanding of basic software development concepts (version control, APIs, data structures)
- Familiarity with command-line interfaces and development environments
- Access to API credits from at least one major LLM provider
- Dedicated study time of 8–15 hours per week

2.1.1 Technical Skills Matrix

Rate your current proficiency (1 = Beginner, 5 = Expert) for each skill:

Skill Area	1	2	3	4	5	Required
Python programming	<input type="checkbox"/>	3+				
REST API consumption	<input type="checkbox"/>	2+				
JSON/YAML data formats	<input type="checkbox"/>	3+				
Git version control	<input type="checkbox"/>	2+				
Database fundamentals (SQL)	<input type="checkbox"/>	2+				
Web development basics	<input type="checkbox"/>	1+				
Machine learning concepts	<input type="checkbox"/>	1+				
Cloud services familiarity	<input type="checkbox"/>	1+				

2.1.2 Prerequisite Gap Remediation

If you identified gaps in the essential prerequisites, the following resources can help you prepare:

Gap Area	Recommended Resource	Time Investment
Python Programming	Python official tutorial; Codecademy Python course	20–40 hours
REST APIs	RESTful API tutorials; Postman learning center	5–10 hours
Git/Version Control	Pro Git book (free online); GitHub Learning Lab	5–10 hours
Command Line	Linux Command Line Basics; terminal tutorials	5–8 hours
JSON/Data Formats	JSON/YAML tutorials; practice with real APIs	2–4 hours
ML Concepts	3Blue1Brown neural networks; fast.ai intro	10–15 hours

2.2 Environment Setup Checklist

Complete this checklist before beginning Week 1:

Self-Evaluation Checklist

Development Environment:

- Python 3.9+ installed with pip/conda
- Code editor or IDE configured (VS Code, PyCharm, etc.)
- Git installed and GitHub account created
- Virtual environment manager available (venv, conda, poetry)
- Jupyter Notebook or JupyterLab installed

API Access:

- Primary LLM provider account and API key (any major provider)
- Secondary provider account (recommended for comparison)
- API credit budget allocated (\$50–100 recommended for full program)
- Environment variables configured for secure key storage

Learning Infrastructure:

- Note-taking system established (Notion, Obsidian, etc.)
- Calendar blocked for study sessions
- Community access (Discord, Slack, or study group)
- Core textbooks acquired (see Resources chapter)

2.3 Mindset Preparation

Success in this program requires embracing several key mindsets:

2.3.1 Empirical Iteration

LLM behavior is inherently probabilistic and sometimes surprising. Effective prompt engineers adopt an experimental mindset, systematically testing hypotheses rather than assuming outcomes. Every prompt is a hypothesis to be validated, not a certainty.

2.3.2 Provider Agnosticism

While you may develop on one platform, the principles taught here transfer across all major providers. Avoid over-fitting your mental models to any single provider's quirks or features. The goal is portable expertise.

2.3.3 Continuous Learning

The LLM landscape evolves rapidly. Techniques that are state-of-the-art today may be superseded within months. This program provides foundational skills and frameworks for ongoing adaptation, not a static body of knowledge.

2.3.4 Production Orientation

From Day 1, think about how techniques would work in production environments. Consider edge cases, failure modes, costs, and user experience even during learning exercises.

Part II

Detailed Learning Path

Chapter 3

Phase 1: Foundation Mastery

⌚ Time Investment


Duration: 4 weeks

Weekly Commitment: 10–12 hours

Total Hours: 40–48 hours

Delivery: Personal Prompt Cookbook

3.1 Phase 1 Learning Path Visualization



3.2 Week 1: Prompt Engineering Fundamentals

3.2.1 Learning Objectives

By the end of Week 1, you will:

- Understand how language models process and generate text at a conceptual level
- Identify the components of effective prompts and their purposes
- Apply fundamental prompting techniques with confidence
- Configure generation parameters appropriately for different tasks
- Recognize the capabilities and limitations of current LLM technology

3.2.2 Day-by-Day Breakdown

Days 1–2: Understanding Language Models

Core Concepts:

- Tokenization: How text becomes numbers; why token boundaries matter
- Context windows: Working within limits; strategies for context management
- Probabilistic generation: Outputs as samples from distributions
- Instruction tuning: How models learn to follow directions (RLHF, Constitutional AI concepts)
- Model families: Understanding capability tiers without vendor lock-in

Practical Activities:

1. Experiment with a tokenizer to understand token boundaries
2. Test identical prompts across different model sizes, noting capability differences
3. Explore generation parameters (temperature, top-p) with the same prompt

Time Investment: 4–5 hours

Days 3–4: Anatomy of Effective Prompts

Core Concepts:

- Prompt structure: Role, context, instruction, format, constraints
- Clarity techniques: Specificity, explicit requirements, unambiguous language
- Output formatting: JSON, XML, Markdown, custom structures
- Positive and negative constraints: What to include versus avoid
- Delimiters and structure: Organizing complex prompts

Practical Activities:

1. Transform 5 vague prompts into precise, well-structured versions
2. Create a “prompt anatomy” template for future use
3. Build a prompt that produces valid JSON output reliably

Time Investment: 4–5 hours

Days 5–7: Core Prompting Techniques

Techniques to Master:

Technique	Implementation Focus
Role Prompting	Assign personas to influence response style, depth, and perspective
Context Setting	Provide relevant background to ground the model's responses
Instruction Clarity	Write unambiguous, specific, actionable instructions
Constraint Specification	Define boundaries: length, scope, format, style
Output Formatting	Specify exact output structure for reliable parsing
Task Decomposition	Break complex requests into sequential steps

Practical Activities:

1. Apply each technique to 3 different task types
2. Compare outputs with and without each technique
3. Document which techniques work best for which scenarios

Time Investment: 3–4 hours

Self-Evaluation Checklist

Week 1 Self-Assessment:

- I can explain why the same prompt might produce different outputs
- I can identify at least 5 components of a well-structured prompt
- I can transform a vague request into a precise prompt
- I can reliably produce structured output (JSON) from prompts
- I understand when to use high vs. low temperature settings
- I have documented at least 6 fundamental prompting techniques

3.3 Week 2: Advanced Prompting Patterns

3.3.1 Learning Objectives

By the end of Week 2, you will:

- Implement chain-of-thought prompting and its variants effectively
- Design few-shot examples that maximize model performance
- Apply self-consistency techniques for high-stakes decisions
- Use critique-and-revise patterns to improve output quality

- Understand when to apply each advanced pattern

3.3.2 Pattern Deep Dives

Chain-of-Thought (CoT) Prompting

Chain-of-thought prompting enables complex reasoning by having the model generate intermediate steps before reaching conclusions.

CoT Variant	Trigger Phrase	Best Use Case
Zero-Shot CoT	“Let’s think step by step”	Quick reasoning enhancement without examples
Few-Shot CoT	Provide reasoning examples	Complex problems requiring specific reasoning style
Self-Consistency	Multiple CoT paths	High-stakes decisions requiring confidence
Tree-of-Thought	Branching exploration	Problems with multiple valid approaches
Program-Aided	Include code execution	Mathematical or logical verification needed

💡 Pro Tip

Chain-of-thought is most effective for multi-step reasoning problems. For simple factual recall or creative writing, CoT may actually reduce quality by introducing unnecessary verbosity. Match the technique to the task.

Few-Shot Learning Mastery

Aspect	Best Practice
Example Count	Start with 3–5 examples; add more only if needed for consistency
Example Selection	Choose diverse, representative cases; include edge cases sparingly
Example Ordering	Place most similar examples closest to the actual task
Format Consistency	Maintain identical structure across all examples
Negative Examples	Include sparingly when boundary clarification is essential
Dynamic Selection	Consider retrieval-based example selection for varied inputs

Iterative Refinement Patterns

Pattern	Implementation	When to Use
Critique-Revise	Generate, self-critique, improve	Quality-critical outputs

Pattern	Implementation	When to Use
Multi-Turn Refinement	Progressive improvement via conversation	Complex, evolving requirements
Self-Evaluation	Assess output against criteria	Before final delivery
Adversarial Testing	Probe for weaknesses	Security-sensitive applications
Constitutional Revision	Apply principle-based evaluation	Alignment-critical outputs

❗ Common Pitfall

Common Pitfall: Over-engineering prompts with too many patterns at once. Start simple and add complexity only when simpler approaches fail. Each additional technique increases token usage and potential failure points.

3.3.3 Week 2 Practical Exercises

Project Deliverable

Exercise 2.1: CoT Implementation Study

Select three distinct problem types:

1. Mathematical word problem
2. Logical reasoning puzzle
3. Code debugging scenario

For each, implement both zero-shot and few-shot CoT. Document:

- Accuracy comparison
- Reasoning quality assessment
- Token usage difference
- Recommendation for when to use each

Deliverable: Comparative analysis document with all prompts and outputs.

Project Deliverable

Exercise 2.2: Few-Shot Engineering

Create a classification task (sentiment, intent, or category assignment):

1. Build few-shot prompts with 1, 3, 5, and 10 examples
2. Test each variant on 20 test cases
3. Measure accuracy, consistency, and token cost
4. Vary example ordering and measure impact

Deliverable: Performance analysis with optimal configuration recommendation.

≡ Self-Evaluation Checklist

Week 2 Self-Assessment:

- I can implement zero-shot and few-shot chain-of-thought prompting
- I can design effective few-shot examples for classification tasks
- I understand when self-consistency improves reliability
- I can implement a critique-and-revise pipeline
- I have documented at least 10 distinct prompting patterns
- I can articulate the trade-offs between different patterns

3.4 Week 3: Systematic Evaluation

3.4.1 Learning Objectives

By the end of Week 3, you will:

- Design appropriate evaluation metrics for different task types
- Create comprehensive test sets including edge cases
- Implement automated evaluation pipelines
- Use LLM-as-judge techniques effectively
- Establish baselines and track improvements systematically

3.4.2 Evaluation Framework

Metric Selection by Task Type

Task Type	Primary Metrics	Implementation Notes
Classification	Accuracy, F1, Precision, Recall	Use stratified test sets; report per-class metrics
Extraction	Exact Match, Token F1	Define matching criteria carefully; handle partial matches
Generation	BLEU, ROUGE, BERTScore	Choose based on what aspect matters (fluency, coverage, similarity)
Summarization	ROUGE-L, factual accuracy	Include human evaluation for faithfulness
Code Generation	Execution pass rate, correctness	Automated testing against test cases
Open-Ended	LLM-as-judge, human rating	Define rubrics explicitly; ensure consistency

Test Set Construction

A high-quality test set includes:

- **Representative Cases (60%)**: Typical inputs the system will encounter
- **Edge Cases (20%)**: Boundary conditions, unusual inputs, extreme values
- **Adversarial Cases (10%)**: Inputs designed to confuse or break the system
- **Regression Cases (10%)**: Previously failed inputs that have been fixed

LLM-as-Judge Implementation

💡 Pro Tip

When using LLMs to evaluate other LLM outputs, provide explicit rubrics with scoring criteria. Include examples of outputs at each score level. Consider using a different model family for evaluation to avoid systematic biases.

3.4.3 Week 3 Practical Exercises

☛ Project Deliverable

Exercise 3.1: Golden Test Set Creation

For a prompt developed in Weeks 1–2:

1. Create 30+ test cases with expected outputs
2. Categorize by case type (representative, edge, adversarial)
3. Implement automated scoring for applicable metrics
4. Run baseline evaluation and document results

Deliverable: Test set file (JSON/YAML) plus evaluation script with baseline results.

☛ Project Deliverable

Exercise 3.2: LLM-as-Judge Pipeline

Build an evaluation system that:

1. Defines clear scoring rubrics (1–5 scale with descriptions)
2. Implements LLM-based evaluation with consistent prompts
3. Calculates inter-rater reliability (if using multiple judge prompts)
4. Generates summary reports with score distributions

Deliverable: Working LLM-as-judge implementation with documentation.

3.5 Week 4: Specialized Prompting Domains

3.5.1 Learning Objectives

By the end of Week 4, you will:

- Apply prompting techniques to code generation tasks effectively
- Produce reliable structured outputs (JSON, XML, YAML)
- Work with vision-language models for multi-modal tasks
- Adapt prompting strategies for domain-specific applications
- Complete the Phase 1 capstone project

3.5.2 Code Generation Patterns

Pattern	Implementation	
Specification-First	Provide detailed requirements including edge cases before requesting code	
Test-Driven	Supply test cases; ask for implementation that passes them	
Incremental Development	Build functionality step-by-step, verifying each step	
Pseudo-code	Scaffold-ing	Outline structure first, then request implementation
Code Review	Prompt-ing	Analyze and improve existing code systematically
Documentation Generation		Generate docstrings, comments, and READMEs from code

3.5.3 Structured Output Reliability

- **Schema Definition:** Always provide explicit JSON schemas or format examples
- **Validation:** Implement parsing with error handling; retry on malformed output
- **Mode Selection:** Use provider-specific structured output modes when available
- **Fallback Strategies:** Define behavior when structured output fails

3.5.4 Multi-Modal Prompting

When working with vision-language models:

- **Image Description:** Be specific about what aspects of the image to analyze
- **Grounding:** Reference image regions explicitly (“in the upper left,” “the red object”)
- **Task Clarity:** Distinguish between description, analysis, and generation tasks
- **Limitations:** Understand model limitations with text in images, fine details, spatial reasoning

3.6 Phase 1 Capstone: Personal Prompt Cookbook

🏁 Milestone

Deliverable: Personal Prompt Cookbook

Compile a comprehensive prompt cookbook containing:

Part 1: Pattern Library (20–30 patterns)

- Pattern name and category
- Description and use cases
- Template with placeholders
- 2–3 concrete examples with outputs
- Known limitations and failure modes
- Provider-specific notes (if any)

Part 2: Evaluation Framework

- Test cases for representative patterns
- Automated evaluation scripts
- Quality metrics and thresholds
- LLM-as-judge rubrics

Part 3: Personal Insights

- Best practices discovered through experimentation
- Anti-patterns to avoid
- Decision framework for pattern selection

Format: GitHub repository with Markdown documentation and supporting code.

⌚ Assessment Checkpoint

Phase 1 Assessment Criteria:

Criterion	Weight
Pattern variety and coverage (20+ distinct patterns)	25%
Quality of documentation and examples	25%
Evaluation implementation and rigor	20%
Personal insights and original contributions	15%
Code quality and organization	15%


Passing Threshold: 70% overall, with no criterion below 50%.

Chapter 4

Phase 2: Application Development

⌚ Time Investment
Duration: 6 weeks
Weekly Commitment: 10–12 hours
Total Hours: 60–72 hours
Delivery: 2–3 Portfolio Applications

4.1 Phase 2 Learning Path Visualization



4.2 Week 5: LLM Application Architecture

4.2.1 Learning Objectives

- Identify and select appropriate architectural patterns for LLM applications
- Design component interactions for maintainability and scalability
- Implement proper state management and conversation handling
- Apply caching strategies for cost and latency optimization
- Handle errors gracefully in LLM-powered systems

4.2.2 Architecture Pattern Catalog

Pattern	Description	Complexity	Use Cases
Direct API	Simple request-response	★ ★ ★ ★	Single-turn tasks, MVP

Pattern	Description	Complexity	Use Cases
Prompt Chain	Sequential prompts	★★★	Multi-step workflows
Router	Classify and delegate	★★★	Multi-intent systems
RAG Pipeline	Retrieve then generate	★★★★	Knowledge-based QA
Agent Loop	Plan-act-observe cycle	★★★★	Autonomous tasks
Multi-Agent	Coordinated specialists	★★★★★	Complex workflows
Human-in-Loop	LLM + human verification	★★★★	High-stakes decisions

4.2.3 Component Design Principles

Prompt Management

- Store prompts in version-controlled templates, not hard-coded strings
- Implement parameterization for dynamic content injection
- Maintain separate prompts for development, testing, and production
- Document prompt purpose, expected inputs, and output format

Context Management

- Implement context window tracking to prevent truncation
- Design summarization strategies for long conversations
- Consider sliding window, summarization, or hybrid approaches
- Handle context overflow gracefully with clear user feedback

Response Handling

- Parse and validate outputs before use
- Implement retry logic with exponential backoff
- Define fallback behaviors for malformed responses
- Log all requests and responses for debugging

Project Deliverable

Exercise 5.1: Architecture Design Document

Design an LLM application (choose: customer support bot, code review assistant, or research summarizer):

1. Create system context and component diagrams
2. Document data flows and API contracts
3. Define error handling and fallback strategies
4. Estimate costs at different usage levels
5. Plan for horizontal scaling

Deliverable: Architecture design document with diagrams.

4.3 Week 6: Tool Integration and Function Calling

4.3.1 Learning Objectives

- Implement function calling across different provider APIs
- Design tool schemas that guide model behavior effectively
- Handle parallel and sequential tool execution
- Manage errors and edge cases in tool chains
- Apply security best practices for tool access

4.3.2 Tool Design Best Practices

Aspect	Best Practice
Naming Convention	Use clear, action-oriented names: <code>search_database</code> , <code>send_email</code> , <code>calculate_total</code>
Description Quality	Provide detailed descriptions of when, why, and how to use each tool
Parameter Design	Well-typed parameters with clear descriptions, constraints, and examples
Return Format	Consistent, parseable return formats; include both data and status
Error Messages	Informative errors the LLM can interpret and respond to appropriately
Scope Limitation	Minimal necessary permissions; explicit boundaries on what tools can do
Confirmation Gates	Require confirmation for destructive or expensive operations

! Common Pitfall

Security Warning: Tools that access external systems (databases, APIs, file systems) must validate all LLM-generated parameters. Never pass LLM output directly to system commands or database queries without sanitization. Implement rate limiting and audit logging for all tool invocations.

Project Deliverable**Project 1: Multi-Tool Agent**

Build an agent with access to 4+ tools:

- Information retrieval tool (web search or database)
- Computation tool (calculator, data analysis)
- Communication tool (email draft, notification)
- File/data tool (read/write structured data)

Requirements:

- Proper tool selection based on user intent
- Sequential and parallel tool execution
- Error handling and recovery
- Conversation history integration
- Security validation on all inputs


Deliverable: Working agent with documentation and test suite.

4.4 Weeks 7–8: Retrieval-Augmented Generation (RAG)

4.4.1 Learning Objectives

- Design and implement complete RAG pipelines
- Apply appropriate document chunking strategies
- Select and configure embedding models and vector stores
- Implement retrieval optimization techniques
- Evaluate and iterate on RAG system quality

4.4.2 RAG Pipeline Architecture



4.4.3 Chunking Strategies

Strategy	Description	Best For
Fixed Size	Split at token/character count	Simple implementation, uniform retrieval
Recursive	Split hierarchically (paragraph, sentence, word)	Varied document structures
Semantic	Split at topic boundaries	Topic-coherent retrieval
Document-Aware	Respect document structure (headers, sections)	Structured documents
Sliding Window	Overlapping chunks	Context preservation
Parent-Child	Small chunks linked to larger context	Precise retrieval with broad context

4.4.4 Retrieval Optimization

Technique	Implementation
Hybrid Search	Combine dense (embedding) and sparse (BM25) retrieval
Query Expansion	Generate multiple query variants to improve recall
HyDE	Generate hypothetical answer, embed it for retrieval
Reranking	Use cross-encoder to reorder initial retrieval results
Metadata Filtering	Pre-filter by date, source, category before semantic search
Multi-Query	Generate diverse queries from single user input

Project Deliverable

Project 2: RAG-Based Document Q&A

Build a complete RAG system:

1. Document ingestion pipeline (PDF, Markdown, HTML)
2. Configurable chunking with multiple strategies
3. Vector storage with any provider-agnostic approach
4. Conversational interface with context tracking
5. Source citation in responses
6. Evaluation suite with retrieval and generation metrics

Optimization Phase:

- Compare chunking strategies on same document set
- Implement and compare reranking approaches
- Test hybrid retrieval vs. dense-only
- Document performance trade-offs

Deliverable: Working RAG application with optimization analysis report.

4.5 Week 9: Conversational Agents and Memory

4.5.1 Learning Objectives

- Design memory architectures for different conversation types
- Implement context summarization and compression
- Build agents with consistent personas across sessions
- Handle entity tracking and reference resolution
- Manage long-term memory persistence

4.5.2 Memory Architecture Comparison

Type	Implementation	Pros	Cons
Buffer	Full history in context	Complete context	Context limit hit quickly
Window	Last N turns only	Predictable tokens	Loses early context
Summary	Compressed history	Handles long chats	Information loss
Entity	Track entities/facts	Efficient storage	Complex implementation

Type	Implementation	Pros	Cons
Vector	Semantic retrieval	Relevant recall	Retrieval quality varies
Hybrid	Combined approaches	Balanced trade-offs	System complexity

4.5.3 Persona Consistency Techniques

- **System Prompt Anchoring:** Define persona characteristics in system prompt
- **Fact Sheets:** Maintain explicit lists of persona facts for reference
- **Consistency Checking:** Validate responses against persona constraints
- **Memory Integration:** Include persona facts in memory retrieval

4.6 Week 10: Integration and Capstone Projects

4.6.1 Portfolio Project Requirements

Complete at least two portfolio-quality applications:

🤖 Track A: Conversational Agent

- Multi-turn chatbot with memory
- Consistent persona across sessions
- Tool integration for actions
- Conversation analytics

💡 Track B: Knowledge System

- Document ingestion pipeline
- Semantic search interface
- Answer with citations
- Evaluation dashboard

📈 Track C: Data Interface

- Natural language to structured queries
- Result visualization
- Query explanation
- Multi-source integration

🏁 Milestone

Phase 2 Capstone Deliverables:

For each project:

- Working application deployed (local or cloud)
- Source code in version control
- Architecture documentation with diagrams
- README with setup and usage instructions
- Evaluation suite with test cases
- Performance metrics and analysis
- Provider abstraction (works with multiple LLMs)

📋 Assessment Checkpoint

Phase 2 Assessment Criteria:

Criterion	Weight
Application functionality and completeness	30%
Code quality, architecture, and maintainability	25%
Documentation quality (README, architecture docs)	15%
Evaluation implementation and coverage	15%
Provider abstraction and portability	15%

Passing Threshold: 70% overall, with working functionality required.

Chapter 5

Phase 3: Production Engineering

⌚ Time Investment


Duration: 4 weeks

Weekly Commitment: 10–12 hours

Total Hours: 40–48 hours

Delivery: Production Playbook

5.1 Phase 3 Learning Path Visualization



5.2 Week 11: System Design for Production

5.2.1 Learning Objectives

- Apply decision frameworks for prompting vs. fine-tuning vs. RAG
- Model costs accurately across different usage scenarios
- Design systems for latency, reliability, and scalability requirements
- Select appropriate models based on comprehensive criteria
- Plan multi-provider strategies for resilience

5.2.2 Decision Framework: Prompting vs. Fine-Tuning vs. RAG

Factor	Better Prompting	Fine-Tuning	RAG	
Knowledge Type	General/static	Style/format	Dynamic/external	
Data Volume	Any	100s–1000s examples	Any corpus size	
Update Frequency	Frequency	Immediate	Requires retraining	Real-time possible
Iteration Speed	Very fast	Days to weeks	Fast (retrieval tuning)	
Cost Structure	Per-token	Training + inference	Storage + retrieval + inference	
Attribution	None	None	Source citation possible	
Hallucination	Model-dependent	Reduced for trained domain	Reduced with good retrieval	

5.2.3 Cost Modeling Framework

- **Input Costs:** System prompt + context + user input tokens
- **Output Costs:** Generated response tokens
- **RAG Costs:** Embedding generation + vector storage + retrieval queries
- **Tool Costs:** External API calls, compute for tool execution
- **Hidden Costs:** Retry attempts, evaluation runs, logging storage

💡 Pro Tip

Cost Optimization Strategies:

1. Cache frequent queries and their responses
2. Use smaller models for classification/routing, larger for generation
3. Compress system prompts; eliminate redundancy
4. Implement early stopping for streaming responses when appropriate
5. Batch requests when real-time response isn't required

5.3 Week 12: Evaluation at Scale

5.3.1 Learning Objectives

- Build comprehensive evaluation pipelines for production systems
- Implement continuous evaluation with automated alerting
- Design and execute A/B tests for prompt changes
- Detect and respond to model drift and regression
- Balance automated and human evaluation appropriately

5.3.2 Multi-Level Evaluation Framework

Level	What to Measure	How to Measure	Frequency
Component	Individual prompt quality	Unit tests, golden sets	Every change
Pipeline	End-to-end correctness	Integration tests, traces	Daily
System	User-facing quality	A/B tests, user feedback	Weekly
Business	Impact on KPIs	Analytics, conversion	Monthly

5.3.3 LLM-as-Judge Best Practices

- **Explicit Rubrics:** Define 1–5 scale with detailed descriptions for each level
- **Calibration Examples:** Include examples of outputs at each score level
- **Multiple Dimensions:** Evaluate relevance, accuracy, completeness, safety separately
- **Judge Diversity:** Use multiple judge prompts or models to reduce bias
- **Human Calibration:** Regularly compare LLM scores to human ratings

Project Deliverable

Exercise 12.1: Production Evaluation Pipeline

Build an evaluation system that:

1. Runs scheduled evaluations against versioned test sets
2. Implements multiple evaluation strategies (exact match, semantic, LLM-judge)
3. Tracks metrics over time with visualization
4. Alerts on regression beyond thresholds
5. Integrates with your Phase 2 application

Deliverable: Working evaluation pipeline with dashboard and alerts.

5.4 Week 13: LLMOps and Production Operations

5.4.1 Learning Objectives

- Implement comprehensive observability for LLM systems
- Design prompt versioning and deployment workflows
- Build feedback collection and processing pipelines
- Establish incident response procedures for LLM failures
- Create continuous improvement loops from production data

5.4.2 Observability Stack Components

Component	Purpose	Key Metrics
Request Logging	Audit trail, debugging	Request/response content, latency, tokens
Distributed Tracing	Multi-step visibility	Span durations, tool calls, retrieval steps
Performance Metrics	System health	Latency percentiles, throughput, error rates
Cost Tracking	Budget management	Token usage, cost per request, cost trends
Quality Monitoring	Output health	Evaluation scores, feedback signals

5.4.3 Prompt Version Management

- **Version Control:** Treat prompts as code; store in Git
- **Change Documentation:** Record rationale for each change
- **Staged Rollout:** Deploy to canary, then percentage, then full
- **Rollback Plan:** Maintain ability to instantly revert
- **A/B Testing Integration:** Support running multiple versions simultaneously

5.5 Week 14: Safety, Security, and Ethics

5.5.1 Learning Objectives

- Identify and mitigate LLM-specific security vulnerabilities
- Implement input validation and output filtering
- Design systems that protect user privacy
- Apply responsible AI principles in practice
- Navigate regulatory requirements (GDPR, AI Act concepts)

5.5.2 Security Vulnerability Catalog

Vulnerability	Description	Mitigation Strategies
Prompt Injection	Malicious input overrides system instructions	Input validation, instruction isolation, output monitoring
Indirect Injection	Malicious content in retrieved documents	Content sanitization, source validation

Vulnerability	Description	Mitigation Strategies
Data Exfiltration	Extracting training data or user data	Output filtering, PII detection, rate limiting
Jailbreaking	Bypassing safety guardrails	Layered defenses, output classification
Excessive Agency	Unintended autonomous actions	Permission boundaries, human-in-loop
Model Extraction	Stealing model via API queries	Rate limiting, query pattern detection

5.5.3 Defense-in-Depth Strategy

1. **Input Layer:** Validate, sanitize, and classify incoming requests
2. **Prompt Layer:** Isolate system instructions from user content
3. **Model Layer:** Use models with safety training; consider guardrails
4. **Output Layer:** Filter, classify, and validate generated content
5. **Monitoring Layer:** Detect anomalies, log for audit, alert on violations

❗ Common Pitfall

Never trust LLM output for security-sensitive operations. Always validate:

- SQL queries before execution
- API calls before making requests
- File paths before access
- Shell commands before running
- Any output that will be rendered as HTML

5.6 Phase 3 Capstone: Production Playbook

Milestone

Deliverable: Production Playbook

Create a comprehensive operational playbook containing:

Part 1: Decision Frameworks

- When to use prompting vs. RAG vs. fine-tuning decision tree
- Model selection criteria and comparison matrix
- Architecture pattern selection guide
- Cost estimation templates and calculators

Part 2: Quality Management

- Evaluation strategy templates by task type
- Test set creation checklists
- LLM-as-judge rubric templates
- Quality metrics definitions and thresholds

Part 3: Operations

- Observability setup guides and dashboards
- Incident response procedures for common failures
- On-call runbooks for LLM-specific issues
- Feedback processing workflows
- Cost monitoring and optimization playbooks

Part 4: Safety and Compliance

- Security hardening checklists
- Input validation and output filtering guidelines
- Privacy protection procedures
- Responsible AI review checklist

Format: Wiki, Notion, or documentation site with templates and checklists.

☒ Assessment Checkpoint**Phase 3 Assessment Criteria:**

Criterion	Weight
Playbook completeness and practicality	30%
Decision framework clarity and usefulness	20%
Operational procedures quality	20%
Security and safety coverage	15%
Production enhancement of Phase 2 project	15%

Passing Threshold: 70% overall.

Chapter 6

Phase 4: Continuous Mastery

⌚ Time Investment

Duration: Ongoing

Weekly Commitment: 3–5 hours

Focus: Continuous professional development

6.1 Mastery Development Framework

Phase 4 represents ongoing professional growth beyond the structured curriculum. Success requires establishing sustainable learning habits and contributing back to the community.

6.1.1 Advanced Technique Areas

Area	Topics to Explore
Multi-Agent Systems	Agent coordination, task decomposition, specialized roles, consensus mechanisms
Advanced Reasoning	Constitutional AI, recursive self-improvement, formal verification of outputs
Efficiency Optimization	Prompt compression, speculative decoding, distillation techniques
Multimodal Integration	Vision-language-action models, embodied AI, cross-modal reasoning
Domain Specialization	Vertical-specific techniques (healthcare, legal, finance, code)
Research Frontiers	Following and implementing academic advances

6.1.2 Specialization Tracks

⚡ Code Generation Specialist

Deep expertise in code generation, completion, review, and debugging. Focus on IDE integration, multi-file context, and repository-level understanding.

Q Knowledge Systems Architect

Advanced RAG systems, knowledge graphs, multi-hop reasoning, and enterprise search. Focus on accuracy, attribution, and scale.

💬 Conversational AI Expert

Advanced dialogue systems, persona consistency, emotional intelligence, and multi-party conversation. Focus on user experience and engagement.

6.1.3 Continuous Learning Practices

- **Weekly:** Read 2–3 new papers or blog posts on LLM techniques
- **Monthly:** Implement one new technique from recent research
- **Quarterly:** Complete a substantial side project exploring new areas
- **Annually:** Contribute to open source, write technical content, or present at meetups

6.1.4 Community Engagement

- Participate in prompt engineering communities (Discord, Reddit, forums)
- Share learnings through blog posts, talks, or tutorials
- Contribute to open-source tools and frameworks
- Mentor others earlier in their learning journey
- Attend conferences and local AI/ML meetups

Part III

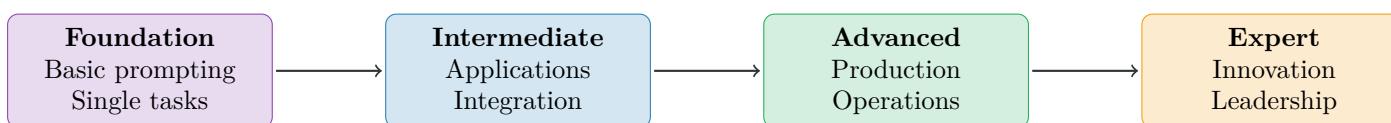
Supporting Materials

Chapter 7

Competency Framework

7.1 Skill Progression Model

This framework defines four competency levels, with specific observable behaviors for each.



7.2 Detailed Competency Matrix

7.2.1 Prompt Engineering Competencies

Skill	Foundation	Intermediate	Advanced	Expert
Basic Prompting	Write clear prompts; get useful outputs	Consistent quality across tasks	Optimize for edge cases	Teach others; create frameworks
Pattern Application	Use 5–10 patterns	Master 20+ patterns; select appropriately	Design custom patterns	Pioneer new patterns
Evaluation	Manual assessment	Automated metrics	Full evaluation pipelines	Design eval frameworks
Multi-Modal	Basic image+text	Complex multi-modal workflows	Production multi-modal systems	Multimodal architecture design

7.2.2 Application Development Competencies

Skill	Foundation	Intermediate	Advanced	Expert
Architecture	Understand patterns	Implement patterns	Design novel architectures	Set organizational standards
Tool Integration	Basic function calls	Multi-tool agents	Complex tool orchestration	Tool framework design

Skill	Foundation	Intermediate	Advanced	Expert
RAG Systems	Basic retrieval	Optimized pipelines	Production at scale	RAG architecture innovation
Agents	Simple agents	Stateful conversational	Multi-agent systems	Agent framework design

7.2.3 Production Engineering Competencies

Skill	Foundation	Intermediate	Advanced	Expert
System Design	Understand trade-offs	Make good decisions	Optimize complex systems	Define best practices
Operations	Basic logging	Full observability	Incident response	Operations strategy
Security	Aware of risks	Implement protections	Comprehensive hardening	Security architecture
Cost Management	Manage costs	Track costs	Optimize costs	Cost-efficient design
				Cost strategy leadership

Chapter 8

Self-Assessment Tools

8.1 Phase Readiness Checklists

8.1.1 Ready for Phase 2?

≡ Self-Evaluation Checklist

Phase 1 Completion Checklist:

Knowledge (must answer “yes” to all):

- Can explain how tokenization affects prompt design
- Can describe at least 5 prompting patterns and when to use them
- Can implement chain-of-thought prompting effectively
- Can design few-shot examples for classification tasks
- Can create evaluation metrics appropriate to different task types

Skills (must answer “yes” to all):

- Have documented 20+ prompt patterns in personal cookbook
- Have implemented automated evaluation for at least one prompt
- Have successfully generated structured output (JSON) reliably
- Have compared prompt performance across different models

Artifacts:

- Personal Prompt Cookbook complete and documented
- Evaluation scripts functional and tested

8.1.2 Ready for Phase 3?

☰ Self-Evaluation Checklist

Phase 2 Completion Checklist: **Knowledge (must answer “yes” to all):**

- Can explain different LLM application architecture patterns
- Can design tool schemas for function calling
- Can describe RAG pipeline components and trade-offs
- Can explain memory architectures for conversational agents
- Can discuss provider abstraction strategies

Skills (must answer “yes” to all):

- Have built working application with tool integration
- Have implemented RAG system with retrieval optimization
- Have built conversational agent with memory
- Applications work with at least two different LLM providers

Artifacts:

- 2–3 portfolio applications complete and documented
- Architecture documentation with diagrams
- Evaluation suites for each application

8.1.3 Ready for Phase 4 (Mastery)?

☰ Self-Evaluation Checklist

Phase 3 Completion Checklist: Knowledge (must answer “yes” to all):

- Can apply decision frameworks for prompting vs. fine-tuning vs. RAG
- Can design comprehensive evaluation pipelines
- Can implement observability for LLM systems
- Can identify and mitigate LLM security vulnerabilities
- Can discuss responsible AI principles and their application

Skills (must answer “yes” to all):

- Have built production evaluation pipeline with alerting
- Have implemented security hardening for LLM application
- Have created operational runbooks for incident response
- Have optimized application for cost and latency

Artifacts:

- Production Playbook complete with decision frameworks
- Phase 2 application enhanced for production
- Operational documentation complete

Chapter 9

Resources and References

9.1 Core Reading List

9.1.1 Primary Textbooks

Title	Focus Area	Best For
<i>Prompt Engineering for Generative AI</i>	Core prompting skills	Phase 1, structured playbook
<i>AI Engineering: Building Applications with Foundation Models</i>	Full lifecycle	Phase 2–3, production focus
<i>LLM Engineer’s Handbook</i>	Operations and infrastructure	Phase 3, production systems
<i>Building LLM Powered Applications</i>	Hands-on development	Phase 2, practical projects
<i>Prompt Engineering for LLMs</i>	Application building	Phase 1–2, comprehensive

9.1.2 Official Documentation (Provider-Agnostic Approach)

Study documentation from multiple providers to understand both common principles and provider-specific features:

- Major LLM provider documentation (prompt engineering guides)
- API references for function calling and tool use
- Best practices and cookbook examples
- Safety and usage policies

9.2 Frameworks and Tools

Category	Examples	Use Case
Orchestration	LangChain, Haystack	LlamaIndex, Building LLM applications
Vector Stores	Pinecone, Weaviate, Chroma, pgvector	RAG systems
Observability	LangSmith, Phoenix, Helicon	Monitoring and tracing
Evaluation	OpenAI Evals, RAGAS, custom	Quality assessment
Safety	NeMo Guardrails, custom filters	Content safety

9.3 Key Research Papers

1. “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models” (Wei et al.)
2. “Self-Consistency Improves Chain of Thought Reasoning” (Wang et al.)
3. “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks” (Lewis et al.)
4. “Constitutional AI: Harmlessness from AI Feedback” (Bai et al.)
5. “Tree of Thoughts: Deliberate Problem Solving with LLMs” (Yao et al.)
6. “ReAct: Synergizing Reasoning and Acting in Language Models” (Yao et al.)
7. “Toolformer: Language Models Can Teach Themselves to Use Tools” (Schick et al.)
8. “The Shift from Models to Compound AI Systems” (Berkeley AI Research)

Chapter 10

Schedule Options and Pacing Guides

10.1 Standard Track (14 Weeks)

Week	Phase	Focus	Hours
1	Phase 1	Fundamentals	10–12
2		Advanced Patterns	10–12
3		Evaluation	10–12
4		Specialization + Capstone	10–12
5	Phase 2	Architecture	10–12
6		Tool Integration	10–12
7		RAG Fundamentals	10–12
8		RAG Optimization	10–12
9		Agents & Memory	10–12
10	Phase 3	Projects + Capstone	10–12
11		System Design	10–12
12		Evaluation at Scale	10–12
13		Operations	10–12
14		Safety + Capstone	10–12

10.2 Accelerated Track (10 Weeks)

For experienced developers with 15–20 hours per week:

Week	Phase	Focus	Hours
1	Phase 1	Fundamentals + Patterns	15–20
2		Evaluation + Specialization + Capstone	15–20
3	Phase 2	Architecture + Tools	15–20
4		RAG Complete	15–20
5		Agents + Memory	15–20
6		Project 1	15–20
7	Phase 3	Projects 2–3	15–20
8		System Design + Evaluation	15–20
9		Operations + Security	15–20
10		Production Playbook	15–20

10.3 Extended Track (18+ Weeks)

For career changers or those with limited weekly time (5–8 hours):

- **Weeks 1–6:** Phase 1 (spread across 6 weeks)
- **Weeks 7–14:** Phase 2 (8 weeks for deeper practice)
- **Weeks 15–18:** Phase 3 (4 weeks at comfortable pace)
- Include additional review weeks as needed

10.4 Self-Paced Guidelines

- **Minimum Commitment:** 5 hours per week to maintain momentum
- **Phase Completion:** Complete each phase before starting the next
- **Project-First Option:** Start projects early, learn concepts as needed
- **Milestone Cadence:** Set weekly goals and track progress
- **Community Support:** Join study groups for accountability

Appendix A

Quick Reference Cards

A.1 Prompting Pattern Quick Reference

Pattern	When to Use	Template Snippet
Role Prompting	Need expert perspective	“You are a [role] with expertise in [domain]...”
Few-Shot	Need consistent format	“Examples: [Ex1] [Ex2] Now: [Task]”
Chain-of-Thought	Complex reasoning	“Let’s think through this step by step...”
Output Format	Structured output needed	“Respond in JSON format: {...}”
Constraints	Specific requirements	“Requirements: 1. ... 2. ... 3. ...”
Self-Consistency	High-stakes decisions	Generate N responses, select most common
Critique-Revise	Quality improvement	“Critique the above, then improve it”
Decomposition	Complex multi-part task	“Break this into steps: [Task]”

A.2 Architecture Decision Quick Reference

Requirement	Recommended Pattern	Key Considerations
Simple Q&A	Direct API	Lowest complexity, highest latency sensitivity
Multi-step workflow	Prompt Chain	Define clear handoffs between steps
Multiple intents	Router Pattern	Classifier accuracy is critical
External knowledge	RAG Pipeline	Chunk size and retrieval quality matter
Tool execution	Agent Loop	Tool design and error handling critical

Requirement	Recommended Pattern	Key Considerations
High stakes	Human-in-Loop	Define clear escalation criteria

A.3 Security Checklist

☰ Self-Evaluation Checklist

Pre-Production Security Checklist:

- Input validation implemented for all user inputs
- System prompt isolated from user content
- Output filtering for PII and sensitive content
- Rate limiting configured
- Audit logging enabled
- Prompt injection test suite passing
- Tool permissions minimized (least privilege)
- Human confirmation for destructive actions
- Error messages don't leak system details
- API keys properly secured (not in code)

Appendix B

Glossary

Term	Definition
Chain-of-Thought (CoT)	Prompting technique that encourages step-by-step reasoning
Context Window	Maximum number of tokens a model can process in a single request
Embedding	Dense vector representation of text for semantic similarity
Few-Shot Learning	Providing examples in the prompt to guide model behavior
Fine-Tuning	Training a model on custom data to adapt its behavior
Function Calling	Enabling models to invoke external functions/tools
Guardrails	Safety mechanisms to constrain model outputs
Hallucination	Model generating plausible but factually incorrect information
HyDE	Hypothetical Document Embedding; generating answer for retrieval
LLM	Large Language Model
LLMOps	Operations practices specific to LLM systems
Prompt Injection	Attack where malicious input overrides system instructions
RAG	Retrieval-Augmented Generation; combining retrieval with generation
Reranking	Re-ordering retrieved results using a more sophisticated model
Self-Consistency	Generating multiple responses and selecting most common answer
System Prompt	Instructions provided to the model that persist across turns
Temperature	Parameter controlling randomness in model outputs
Token	Basic unit of text processing in LLMs
Top-p (Nucleus)	Sampling parameter limiting to most probable tokens
Vector Store	Database optimized for storing and querying embeddings

Term	Definition
Zero-Shot	Prompting without providing examples